

“Don't expect to be easy to work, because if you're easy to do, your heart will be flattered and arrogant.”
_ Buddha



APT: PDE_Machine Learning

Presenter: Dr. Ha Viet Uyen Synh.

Warm-up: Face Recognition

HOW 2D FACIAL SCANNERS RECORD IDENTITIES

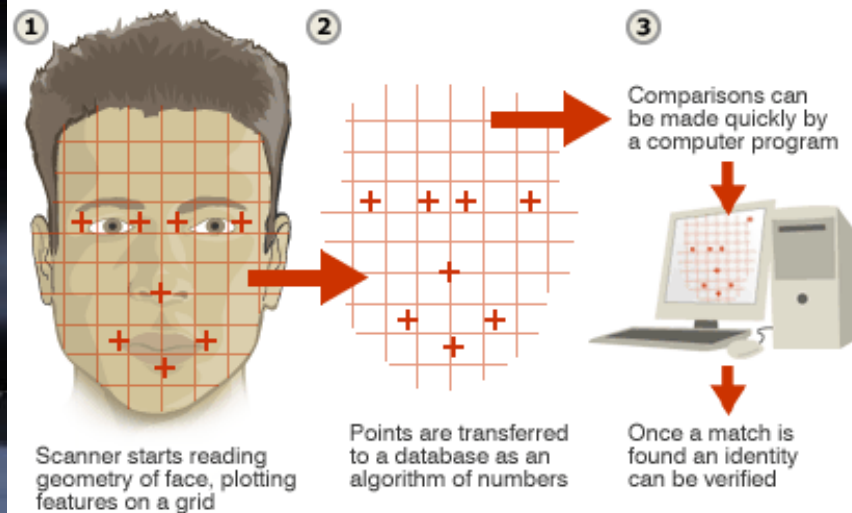



Image Segmentation



Warm-up: U-Net

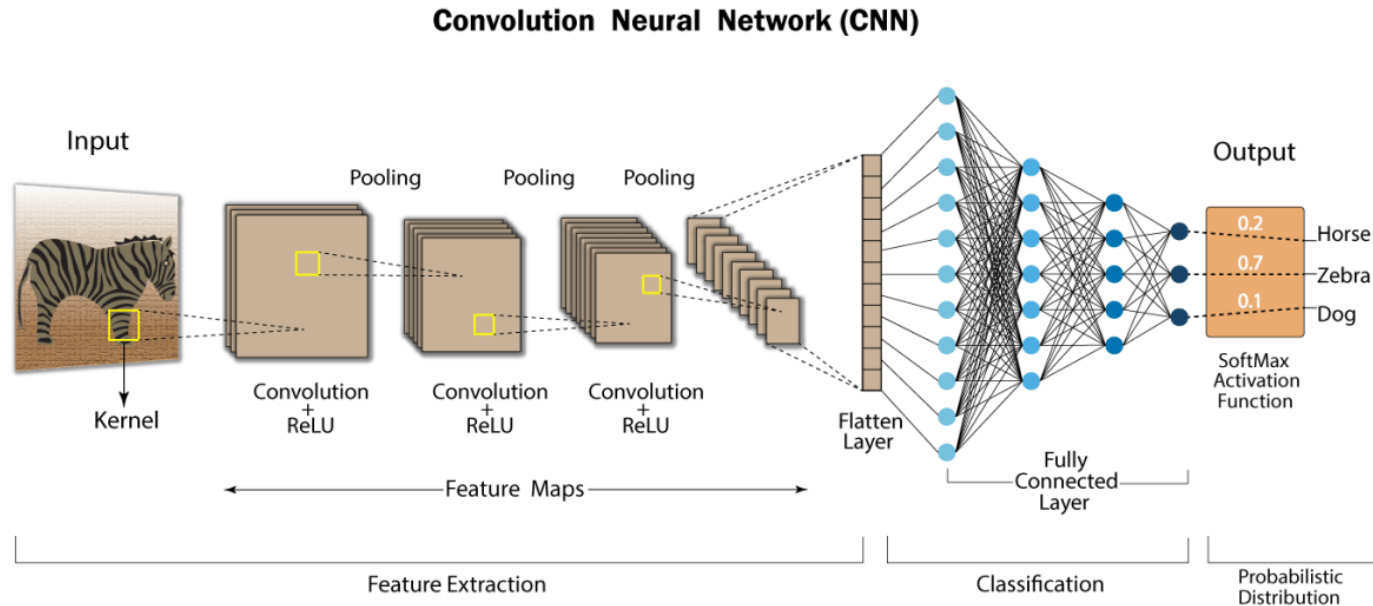




Part A

CONVOLUTIONAL NEURAL NETWORK

1. Convolutional Neural Network _ CNN



In CNN, each input image will pass through a sequence of convolution layers along with pooling, fully connected layers, filters (Also known as kernels). After that, we will apply the Soft-max function to classify an object with probabilistic values 0 and 1.


Convolution Layer

Convolution layer is the first layer to extract features from an input image. By learning image features using a small square of input data, the convolutional layer preserves the relationship between pixels. It is a mathematical operation which takes two inputs such as image matrix and a kernel or filter.

The dimension of the image matrix is $h \times w \times d$.

The dimension of the filter is $f_h \times f_w \times d$.

The dimension of the output is $(h - [f_h/2] + 1) \times (w - [f_w/2] + 1) \times 1$.



Ex:

Image matrix multiplies kernel or filter matrix

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \end{bmatrix} \times \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 4 & 3 & 4 \\ 2 & 4 & 3 \\ 2 & 3 & 4 \end{bmatrix}$$

Convolved Feature

Strides (recall)

1	2	3	4	5	6	7
11	12	13	14	15	16	17
21	22	23	24	25	26	27
31	32	33	34	35	36	37
41	42	43	44	45	46	47
51	52	54	55	56	57	17
61	62	63	64	65	66	67
71	72	73	74	75	76	77

Strides


Convolve with 3*3
filters filled with ones

108	126	
288	306	


Padding (recall)

One time
Cover Pixel

Padding




*



=

One time
Cover Pixel

Pooling Layer



Max Pooling

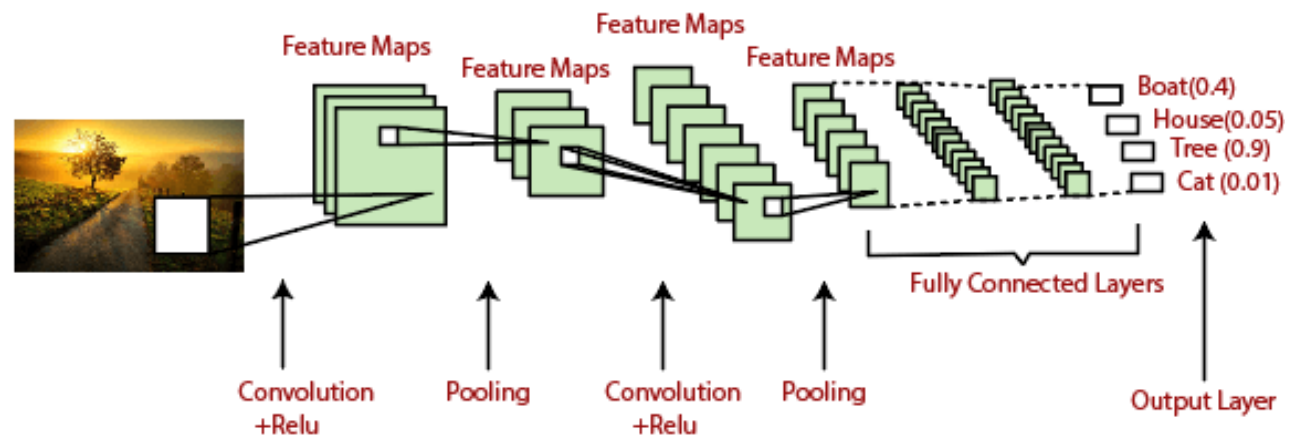

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	2

$\xrightarrow{2 \times 2 \text{ Max-Pool}}$

20	30
112	372

Fully Connected Layer

Fully Connected Layer





2. Image Transforms

Step 1:

```
import torch
import matplotlib.pyplot as plt
import numpy as np
from torchvision import datasets, transforms
```

Step 2:

```
transform1=transforms.Compose([transforms.ToTensor(), tran
sforms.Normalize((0.5,), (0.5,))])
training_dataset=datasets.MNIST(root='./data', train=True,
download=True, transform=transform1)
training_loader=torch.utils.data.DataLoader(dataset=train
ing_dataset, batch_size=100, shuffle=True)
```


Note: Our MNIST images are 28*28 grayscale images which would imply that each image is a two dimensional number by array 28 pixels wide and 28 pixels long and each pixel intensity ranging from 0 to 255.

Step 3:

```
def im_convert(tensor):
    image=tensor.clone().detach().numpy()
    image=image.transpose(1,2,0)
    print(image.shape)
    image=image*(np.array((0.5,0.5,0.5))+np.array((0.5,0.5,0.5)))
    image=image.clip(0,1)
    return image
```

Step 4:

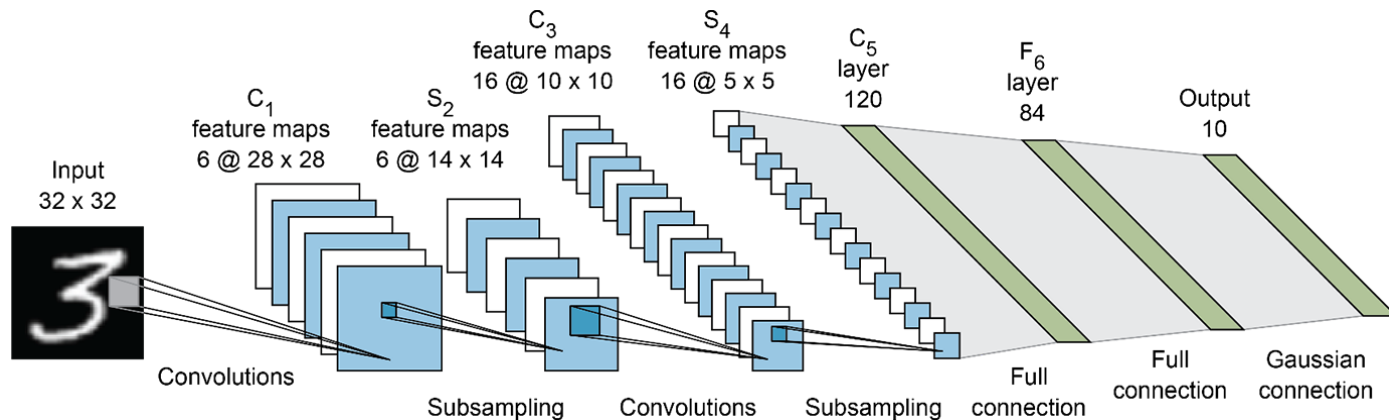
```
dataiter=iter(training_loader)
images,labels=dataiter.next()
fig=plt.figure(figsize=(25,4))
for idx in np.arange(20):
    ax=fig.add_subplot(2,10,idx+1)
    plt.imshow(im_convert(images[idx]))
    ax.set_title([labels[idx].item()])
plt.show()
```



3. CNN Implementation

```
class LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1=nn.Conv2d(1,20,5,1)
        self.conv2=nn.Conv2d(20,50,5,1)
        self.fully1=nn.Linear(4*4*50,500)
        self.fully2=nn.Linear(500,10)

    def forward(self,x):
        x=func.relu(self.conv1(x))
        x=func.max_pool2d(x,2,2)
        x=func.relu(self.conv2(x))
        x=func.max_pool2d(x,2,2)
        x=x.view(-1,4*4*50)
        x=func.relu(self.fully1(x))
        x=self.fully2(x)
        return x
```





4. Training of CNN

Step 1:

```
device=torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
transform1=transforms.Compose([transforms.Resize((28,28)), transforms.ToTensor(), transforms.Normalize((0.5,), (0.5,))])
```

```
training_dataset=datasets.MNIST(root='./data', train=True, download=True, transform=transform1)
```

```
validation_dataset=datasets.MNIST(root='./data', train=False, download=True, transform=transform1)
```

```
training_loader=torch.utils.data.DataLoader(dataset=training_dataset, batch_size=100, shuffle=True)
```

```
validation_loader=torch.utils.data.DataLoader(dataset=validation_dataset, batch_size=100, shuffle=False)
```

Step 2:

```
model=LeNet().to(device)
```


```
criterion=nn.CrossEntropyLoss()
```

```
optimizer=torch.optim.Adam(model.parameters(), lr=0.00001)
```




Step 3:

```
epochs=12
loss_history=[]
correct_history=[]
val_loss_history=[]
val_correct_history=[]
for e in range(epochs):
    loss=0.0
    correct=0.0
    val_loss=0.0
    val_correct=0.0
    for input,labels in training_loader:
        input=input.to(device)    #device
        labels=labels.to(device)  #device
        outputs=model(input)
        loss1=criterion(outputs,labels)
        optimizer.zero_grad()
        loss1.backward()
        optimizer.step()
        _,preds=torch.max(outputs,1)
        loss+=loss1.item()
    correct+=torch.sum(preds==labels.data)
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked in a slightly offset manner, creating a sense of balance and harmony. The background is a soft, out-of-focus landscape with a light sky and a calm body of water.

```
else:
    with torch.no_grad():
        for val_input, val_labels in validation_loader
:
            val_input=val_input.to(device)
            val_labels=val_labels.to(device)
            val_outputs=model(val_input)
            val_loss1=criterion(val_outputs, val_labels
)
            _, val_preds=torch.max(val_outputs, 1)
            val_loss+=val_loss1.item()
            val_correct+=torch.sum(val_preds==val_lab
els.data)

    epoch_loss=loss/len(training_loader)
    epoch_acc=correct.float()/len(training_loader)
    loss_history.append(epoch_loss)
    correct_history.append(epoch_acc)
    val_epoch_loss=val_loss/len(validation_loader)
    val_epoch_acc=val_correct.float()/len(validation_
loader)

    val_loss_history.append(val_epoch_loss)
    val_correct_history.append(val_epoch_acc)
    print('training_loss:{:.4f},{:.4f}'.format(epoch_
loss, epoch_acc.item()))
    print('validation_loss:{:.4f},{:.4f}'.format(val_
epoch_loss, val_epoch_acc.item()))
```

A stack of five smooth, dark, rounded stones is positioned on the left side of the image. They are stacked vertically, with the largest stone at the bottom and the smallest at the top. The stones are resting on a highly reflective surface, which creates a clear mirror image of the stack below it. The background is a soft, out-of-focus light blue and white, suggesting a calm, possibly aquatic or misty environment.

Result:



```
training_loss:1.8134,63.1500
validation_loss:1.0848,79.7000
training_loss:0.7461,83.5250
validation_loss:0.5079,86.8700
training_loss:0.4459,88.2317
validation_loss:0.3603,90.1300
training_loss:0.3410,90.6333
validation_loss:0.2901,91.9600
training_loss:0.2827,92.0683
validation_loss:0.2442,93.1200
training_loss:0.2431,93.0583
validation_loss:0.2127,93.8600
training_loss:0.2138,93.8633
validation_loss:0.1863,94.5200
training_loss:0.1906,94.5000
validation_loss:0.1690,94.9800
training_loss:0.1719,95.0100
validation_loss:0.1509,95.5300
training_loss:0.1562,95.4267
validation_loss:0.1376,95.9000
training_loss:0.1431,95.8233
validation_loss:0.1273,96.2200
training_loss:0.1323,96.1267
validation_loss:0.1164,96.4900
```

Step 4:

```
plt.plot(loss_history, label='Training Loss')
plt.plot(val_loss_history, label='Validation Loss')
plt.legend()
plt.show()

plt.plot(correct_history, label='Training accuracy')
plt.plot(val_correct_history, label='Validation accuracy')

plt.legend()
plt.show()
```





5. Validation of CNN

Step 1:

```
import torch
import matplotlib.pyplot as plt
import numpy as np
import torch.nn.functional as func
import PIL.ImageOps
from torch import nn
from torchvision import datasets, transforms
```

Step 2:

```
device=torch.device("cuda:0" if torch.cuda.is_available()
    else "cpu")
transform1=transforms.Compose([transforms.Resize((28,28))
    ,transforms.ToTensor(),transforms.Normalize((0.5,),(0.5,))
])
training_dataset=datasets.MNIST(root='./data',train=True,
    download=True,transform=transform1)
validation_dataset=datasets.MNIST(root='./data',train=False,
    download=True,transform=transform1)
training_loader=torch.utils.data.DataLoader(dataset=train
    ing_dataset,batch_size=100,shuffle=True)
validation_loader=torch.utils.data.DataLoader(dataset=val
    idation_dataset,batch_size=100,shuffle=False)
```

Step 3:


```
class LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1=nn.Conv2d(1,20,5,1)
        self.conv2=nn.Conv2d(20,50,5,1)
        self.fully1=nn.Linear(4*4*50,500)
        self.dropout1=nn.Dropout(0.5) #<=====
        self.fully2=nn.Linear(500,10)
    def forward(self,x):
        x=func.relu(self.conv1(x))
        x=func.max_pool2d(x,2,2)
        x=func.relu(self.conv2(x))
        x=func.max_pool2d(x,2,2)
        x=x.view(-1,4*4*50)
        x=func.relu(self.fully1(x))
        x=self.dropout1(x) #<=====
        x=self.fully2(x) #<=====
        return x
```

Step 4:


```
model=LeNet().to(device)
criterion=nn.CrossEntropyLoss()
optimizer=torch.optim.Adam(model.parameters(),lr=0.00001)
```

Dropout Layer


Dropout Layer is a technique to reduce overfitting occurs



Dropout Layer



Standard Neural Network




After Applying Dropout

Dropout layers will be placed between the convolutional layers and between the fully connected layers. The dropout layer is used in between layers which have a high number of parameters because these high parameter layers are more likely to overfit and memorize the training data.

Step 5:

```
epochs=12
loss_history=[]
correct_history=[]
val_loss_history=[]
val_correct_history=[]
for e in range(epochs):
    loss=0.0
    correct=0.0
    val_loss=0.0
    val_correct=0.0
    for input,labels in training_loader:
        input=input.to(device)    #<====
        labels=labels.to(device)  #<====
        outputs=model(input)
        loss1=criterion(outputs,labels)
        optimizer.zero_grad()
        loss1.backward()
        optimizer.step()
        _,preds=torch.max(outputs,1)
        loss+=loss1.item()
        correct+=torch.sum(preds==labels.data)
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked horizontally, and their reflection is visible in the water below. The background is a soft, out-of-focus landscape with a light sky and water.

```
else:
    with torch.no_grad():
        for val_input, val_labels in validation_loader:

            val_input=val_input.to(device)
            val_labels=val_labels.to(device)
            val_outputs=model(val_input)
            val_loss1=criterion(val_outputs, val_labels)

            _, val_preds=torch.max(val_outputs, 1)
            val_loss+=val_loss1.item()
            val_correct+=torch.sum(val_preds==val_labels.data)



        epoch_loss=loss/len(training_loader)
        epoch_acc=correct.float()/len(training_loader)
        loss_history.append(epoch_loss)
        correct_history.append(epoch_acc)
        val_epoch_loss=val_loss/len(validation_loader)
        val_epoch_acc=val_correct.float()/len(validation_loader)

        val_loss_history.append(val_epoch_loss)
        val_correct_history.append(val_epoch_acc)
        print('training_loss:{:.4f},{:.4f}'.format(epoch_loss, epoch_acc.item()))
        print('validation_loss:{:.4f},{:.4f}'.format(val_epoch_loss, val_epoch_acc.item()))
```

Step 6:

```
plt.plot(loss_history, label='Training Loss')
plt.plot(val_loss_history, label='Validation Loss')
plt.legend()
plt.show()

plt.plot(correct_history, label='Training accuracy')
plt.plot(val_correct_history, label='Validation accuracy')
plt.legend()
plt.show()
```





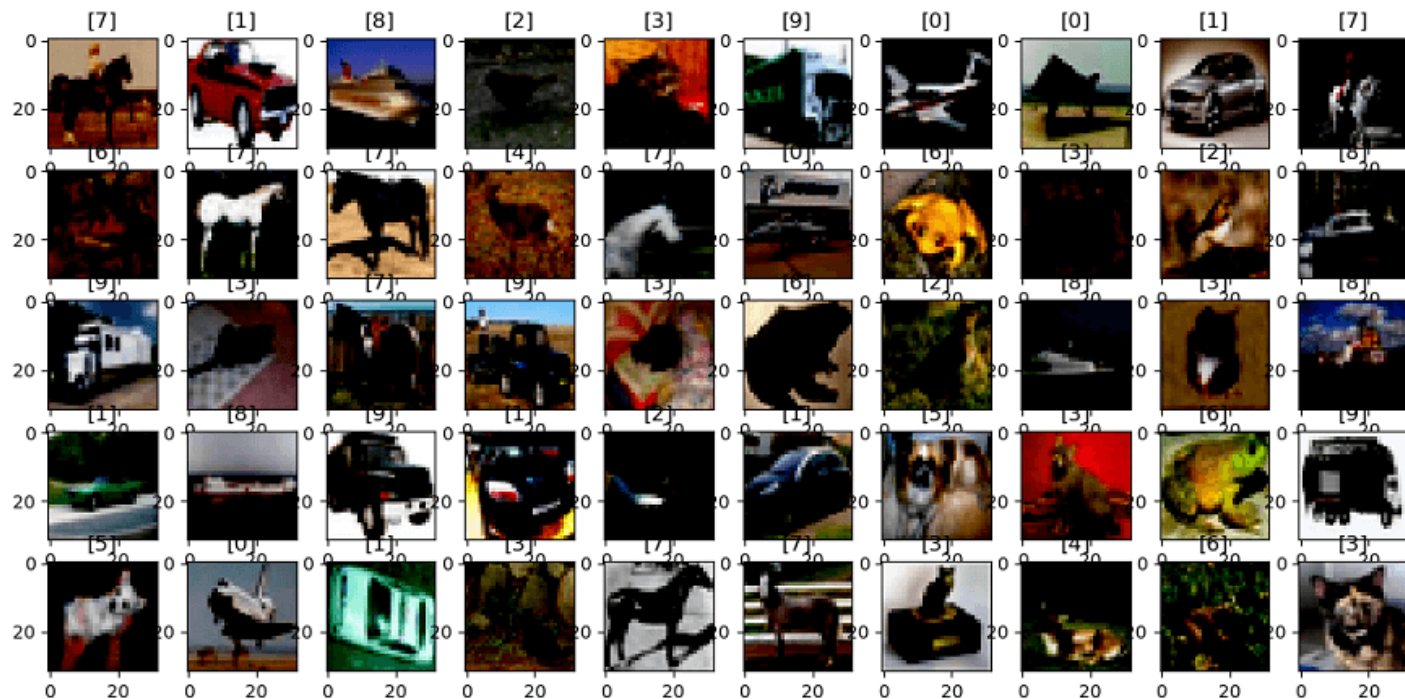
Part B

IMAGE CLASSIFICATION

CIFAR-10 Dataset

The **CIFAR 10**(**C**anadian **I**nstitute for **A**dvanced **R**esearch) will be harder to classify and will come with new barriers which we will need to overcome. It is a collection of the image which is commonly used to train machine learning and computer vision algorithms.

The CIFAR-10 dataset consists of 60000 thirty by thirty color images in 10 classes means 6000 images per class. This dataset is divided into one test batch and five training batches. Every batch contains 10000 images. In the test batch, there are 1000 images which are randomly selected from each class.



CIFAR-100

Dataset



S. No	Superclass	Classes
1.	aquatic mammals	beaver, dolphin, otter, seal, whale
2.	flowers	orchids, poppies, roses, sunflowers, tulips
3.	fish	aquarium fish, flatfish, ray, shark, trout
4.	food containers	bottles, bowls, cans, cups, plates
5.	household electrical devices	clock, computer keyboard, lamp, telephone, television
6.	fruit and vegetables	apples, mushrooms, oranges, pears, sweet peppers
7.	household furniture	bed, chair, couch, table, wardrobe
8.	large carnivores	bear, leopard, lion, tiger, wolf
9.	insects bee, beetle, butterfly, caterpillar, cockroach	
10.	large man-made outdoor things	bridge, castle, house, road, skyscraper
11.	large natural outdoor scenes	cloud, forest, mountain, plain, sea
12.	medium-sized mammals	fox, porcupine, possum, raccoon, skunk
13.	large omnivores and herbivores	camel, cattle, chimpanzee, elephant, kangaroo
14.	non-insect invertebrates	crab, lobster, snail, spider, worm
15.	reptiles	crocodile, dinosaur, lizard, snake, turtle
16.	people	baby, boy, girl, man, woman
17.	trees	maple, oak, palm, pine, willow
18.	small mammals	hamster, mouse, rabbit, shrew, squirrel
19.	vehicles 1	bicycle, bus, motorcycle, pickup truck, train
20.	vehicles 2	lawn-mower, rocket, streetcar, tank, tractor



1. LeNet Model for CIFAR-10 Dataset

Step 1:

```
import torch
import matplotlib.pyplot as plt
import numpy as np
import torch.nn.functional as func
import PIL.ImageOps
from torch import nn
from torchvision import datasets, transforms
import requests
from PIL import Image
```

Step 2:

```
device=torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
transform1=transforms.Compose([transforms.Resize((32, 32)), transforms.ToTensor(), transforms.Normalize((0.5, ), (0.5, ))])
```

```
training_dataset=datasets.CIFAR10(root='./data', train=True, download=True, transform=transform1)
```

```
validation_dataset=datasets.CIFAR10(root='./data', train=False, download=True, transform=transform1)
```

```
training_loader=torch.utils.data.DataLoader(dataset=training_dataset, batch_size=100, shuffle=True)
```

```
validation_loader=torch.utils.data.DataLoader(dataset=validation_dataset, batch_size=100, shuffle=False)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz

170499072/? [00:03<00:00, 52671694.25it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data

Files already downloaded and verified

Step 3:

```
def im_convert(tensor):
    image=tensor.cpu().clone().detach().numpy()
    image=image.transpose(1,2,0)
    print(image.shape)
    image=image*(np.array((0.5,0.5,0.5))+np.array((0.5,0.5,0.5)))
    image=image.clip(0,1)
    return image
```

```
classes=('plane','car','bird','cat','deer','dog','frog','horse','ship','truck')
```

```
dataiter=iter(training_loader)
```

```
images,labels=dataiter.next()
```

```
fig=plt.figure(figsize=(25,4))
```


```
for idx in np.arange(20):
```

```
    ax=fig.add_subplot(2,10,idx+1)
```

```
    plt.imshow(im_convert(images[idx]))
```

```
    ax.set_title(classes[labels[idx].item()])
```

```
(32, 32, 3)
(32, 32, 3)
(32, 32, 3)
(32, 32, 3)
(32, 32, 3)
(32, 32, 3)
(32, 32, 3)
```



Step 4:

```
class LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1=nn.Conv2d(3,20,5,1)
        self.conv2=nn.Conv2d(20,50,5,1)
        self.fully1=nn.Linear(5*5*50,500)
        self.dropout1=nn.Dropout(0.5)
        self.fully2=nn.Linear(500,10)
    def forward(self,x):
        x=func.relu(self.conv1(x))
        x=func.max_pool2d(x,2,2)
        x=func.relu(self.conv2(x))
        x=func.max_pool2d(x,2,2)
        x=x.view(-1,5*5*50)
        x=func.relu(self.fully1(x))
        x=self.dropout1(x)
        x=self.fully2(x)
        return x

model=LeNet().to(device)
criterion=nn.CrossEntropyLoss()
optimizer=torch.optim.Adam(model.parameters(),lr=0.00001)
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked horizontally, and their reflection is visible in the water below. The background is a soft, out-of-focus blue and white gradient.

Step 5:

```
epochs=12
loss_history=[]
correct_history=[]
val_loss_history=[]
val_correct_history=[]
for e in range(epochs):
    loss=0.0
    correct=0.0
    val_loss=0.0
    val_correct=0.0
    for input,labels in training_loader:
        input=input.to(device)
        labels=labels.to(device)
        outputs=model(input)
        loss1=criterion(outputs,labels)
        optimizer.zero_grad()
        loss1.backward()
        optimizer.step()
        _,preds=torch.max(outputs,1)
        loss+=loss1.item()
        correct+=torch.sum(preds==labels.data)
```

A stack of five smooth, dark, rounded stones is placed on a highly reflective surface, likely water. The stones are stacked horizontally, with the top stone being the most prominent. The surface reflects the stones and the background, creating a clear mirror image. The background is a soft, out-of-focus blue and white, suggesting a sky or water surface.

```
else:
    with torch.no_grad():
        for val_input, val_labels in validation_loader
        :
            val_input=val_input.to(device)
            val_labels=val_labels.to(device)
            val_outputs=model(val_input)
            val_loss1=criterion(val_outputs, val_labels
            )
            _, val_preds=torch.max(val_outputs, 1)
            val_loss+=val_loss1.item()
            val_correct+=torch.sum(val_preds==val_lab
            els.data)
            epoch_loss=loss/len(training_loader)
            epoch_acc=correct.float()/len(training_loader)
            loss_history.append(epoch_loss)
            correct_history.append(epoch_acc)
            val_epoch_loss=val_loss/len(validation_loader)
            val_epoch_acc=val_correct.float()/len(validation_
            loader)
            val_loss_history.append(val_epoch_loss)
            val_correct_history.append(val_epoch_acc)
            print('training_loss:{:.4f},{:.4f}'.format(epoch_
            loss, epoch_acc.item()))
            print('validation_loss:{:.4f},{:.4f}'.format(val_
            epoch_loss, val_epoch_acc.item()))
```





```
training_loss:2.2006,21.3820
validation_loss:2.0616,27.2700
training_loss:1.9963,28.2420
validation_loss:1.9451,30.1700
training_loss:1.9216,30.7620
validation_loss:1.8916,32.2000
training_loss:1.8724,32.7180
validation_loss:1.8410,33.9000
training_loss:1.8256,34.4120
validation_loss:1.7974,35.3800
training_loss:1.7788,36.3240
validation_loss:1.7500,37.1300
training_loss:1.7393,37.5580
validation_loss:1.7179,37.5900
training_loss:1.7018,38.8580
validation_loss:1.6762,39.6100
training_loss:1.6696,40.2460
validation_loss:1.6521,40.0500
training_loss:1.6407,40.7400
validation_loss:1.6280,41.1300
training_loss:1.6163,41.8280
validation_loss:1.6038,42.2600
training_loss:1.5954,42.4840
validation_loss:1.5867,43.3300
```


Step 6:

```
plt.plot(loss_history, label='Training Loss')
plt.plot(val_loss_history, label='Validation Loss')
plt.legend()
plt.show()

plt.plot(correct_history, label='Training accuracy')
plt.plot(val_correct_history, label='Validation accuracy')

plt.legend()
plt.show()
```

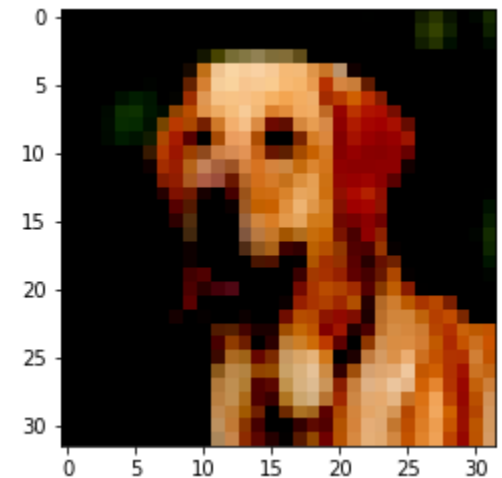




Step 7:

```
url='https://3c1703fe8d.site.internapcdn.net/newman/gfx/news/hires/2018/2-dog.jpg'
response=requests.get(url,stream=True)
img=Image.open(response.raw)
img=transform1(img)
image1=img.to(device).unsqueeze(0)
output=model(image1)
_,pred=torch.max(output,1)
print(classes[pred.item()])
```

dog

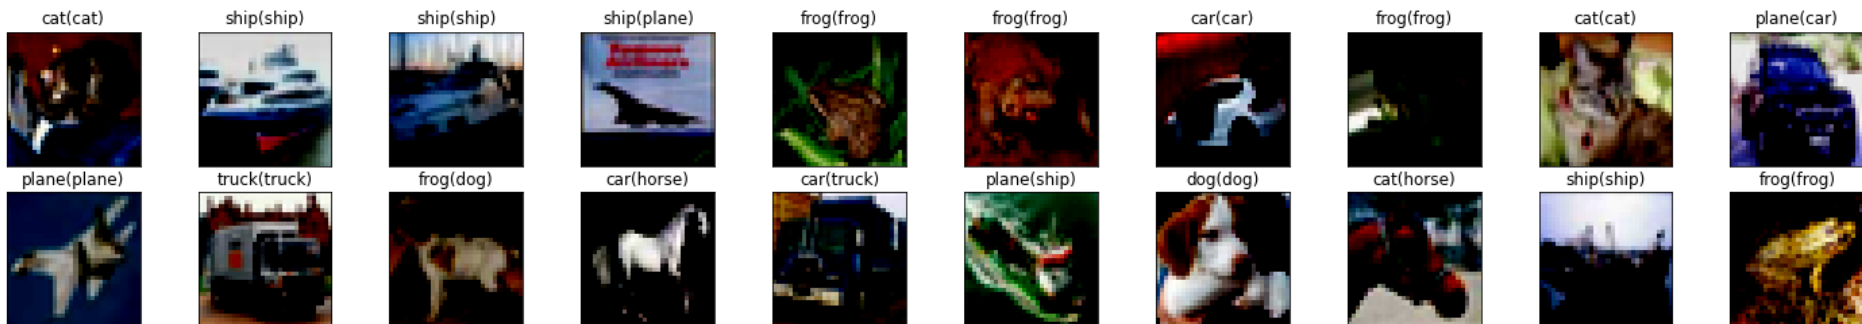





Step 8:

```
dataiter=iter(validation_loader)
images,labels=dataiter.next()
images_=images.to(device)
labels=labels.to(device)
output=model(images_)
_,preds=torch.max(output,1)
fig=plt.figure(figsize=(25,4))
for idx in np.arange(20):
    ax=fig.add_subplot(2,10,idx+1,xticks=[],yticks=[])
    plt.imshow(im_convert(images[idx]))
    ax.set_title("{} ({} )".format(str(classes[preds[idx].item()]),str(classes[labels[idx].item()]),color=("green" if preds[idx]==labels[idx] else "red"))
plt.show()
```

(32, 32, 3)
(32, 32, 3)
(32, 32, 3)
(32, 32, 3)
(32, 32, 3)







3. Hyperparameter Tuning Technique



The fine-tuning of the model is important and can improve the model performance significantly

1. The first modification will be focused on the learning rate.
2. The second modification is very effective. We will simply add more convolution of layers.
3. A larger kernel implies more parameters. We will use a smaller kernel size to remove overfitting.



Old model
Lr =0.00001

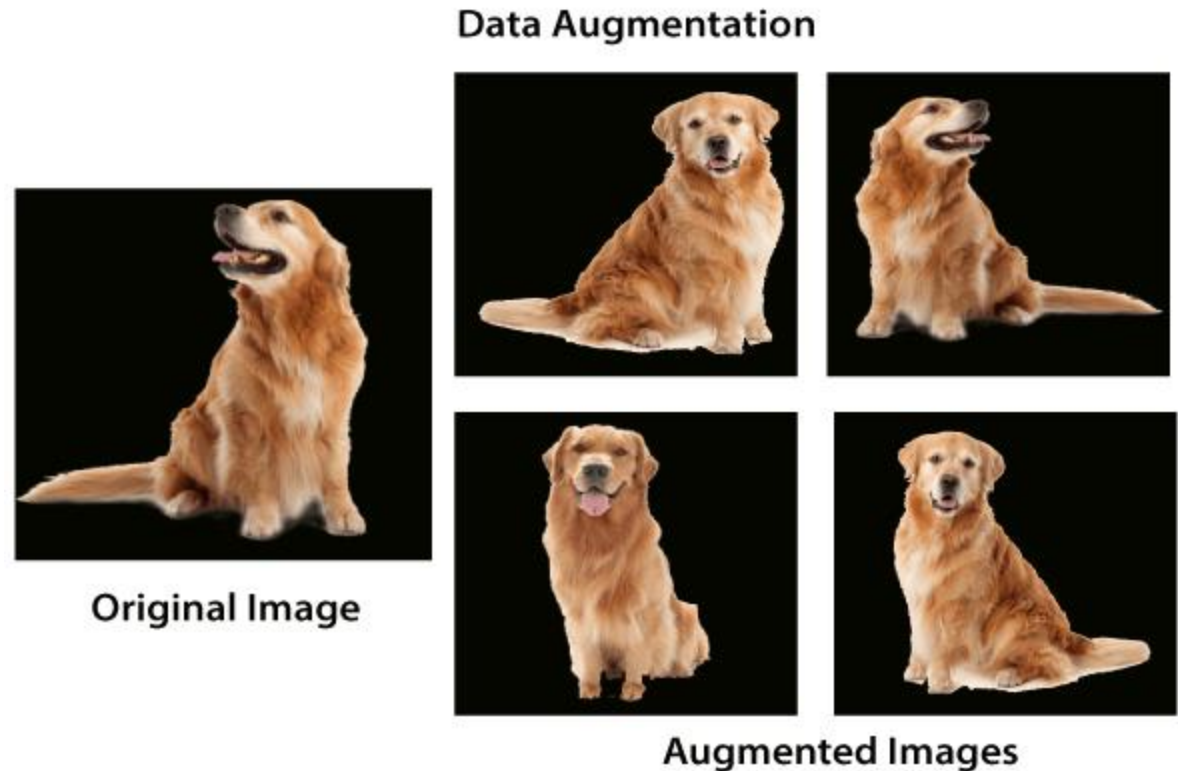
New model
Lr =0.001



```
class LeNet(nn.Module): #Modify
    def __init__(self):
        super().__init__()
        #self.conv1=nn.Conv2d(3,20,5,1)
        #self.conv2=nn.Conv2d(20,50,5,1)
        self.conv1=nn.Conv2d(3,16,5,1)
        self.conv2=nn.Conv2d(16,32,5,1)
        self.conv3=nn.Conv2d(32,64,5,1)
        #self.fully1=nn.Linear(5*5*50,500)
        self.fully1=nn.Linear(5*5*64,500)
        self.dropout1=nn.Dropout(0.5)
        self.fully2=nn.Linear(500,10)
    def forward(self,x):
        x=func.relu(self.conv1(x))
        x=func.max_pool2d(x,2,2)
        x=func.relu(self.conv2(x))
        x=func.max_pool2d(x,2,2)
        #x=x.view(-1,5*5*50)
        x=x.view(-1,5*5*64)
        x=func.relu(self.fully1(x))
        x=self.dropout1(x)
        x=self.fully2(x)
        return x
```

```
model=LeNet().to(device)
criterion=nn.CrossEntropyLoss()
optimizer=torch.optim.Adam(model.parameters(),lr=0.001)
```

4. Data Augmentation Process



The data augmentation technique is useful because it allows our model to look at each image in our dataset from a variety of different perspective. After applying the transformation, the newly created images are known as augmented images because they essentially allow us to augment our dataset by adding new data to it.



Data Augmentation Implementation

Step 1:

```
import torch
import matplotlib.pyplot as plt
import numpy as np
import torch.nn.functional as func
import PIL.ImageOps
from torch import nn
from torchvision import datasets, transforms
import requests
from PIL import Image
```




Step 2:

```
device=torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
transform_train=transforms.Compose([transforms.Resize((32,32)),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(10),
    transforms.RandomAffine(0,shear=10,scale=(0.8,1.2)),
    transforms.ColorJitter(brightness=0.2,contrast=0.2,saturation=0.2),
    transforms.ToTensor(),
    transforms.Normalize((0.5,),(0.5,))])
```

```
transform1=transforms.Compose([transforms.Resize((32,32)),transforms.ToTensor(),transforms.Normalize((0.5,),(0.5,))])
```

```
training_dataset=datasets.CIFAR10(root='./data',train=True,download=True,transform=transform_train)
```

```
validation_dataset=datasets.CIFAR10(root='./data',train=False,download=True,transform=transform1)
```

```
training_loader=torch.utils.data.DataLoader(dataset=training_dataset,batch_size=100,shuffle=True)
```

```
validation_loader=torch.utils.data.DataLoader(dataset=validation_dataset,batch_size=100,shuffle=False)
```

Step 3:

```
def im_convert(tensor):
    image=tensor.cpu().clone().detach().numpy()
    image=image.transpose(1,2,0)
    print(image.shape)
    image=image*(np.array((0.5,0.5,0.5))+np.array((0.
5,0.5,0.5)))
    image=image.clip(0,1)
    return image

classes=('plane','car','bird','cat','deer','dog','fro
g','horse','ship','truck')

dataiter=iter(training_loader)
images,labels=dataiter.next()
fig=plt.figure(figsize=(25,4))
for idx in np.arange(20):
    ax=fig.add_subplot(2,10,idx+1)
    plt.imshow(im_convert(images[idx]))
    ax.set_title(classes[labels[idx].item()])
```



Step 4:

```
class LeNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1=nn.Conv2d(3,16,3,1, padding=1)
        self.conv2=nn.Conv2d(16,32,3,1, padding=1)
        self.conv3=nn.Conv2d(32,64,3,1, padding=1)
        self.fully1=nn.Linear(4*4*64,500)
        self.dropout1=nn.Dropout(0.5)
        self.fully2=nn.Linear(500,10)


    def forward(self,x):
        x=func.relu(self.conv1(x))
        x=func.max_pool2d(x,2,2)
        x=func.relu(self.conv2(x))
        x=func.max_pool2d(x,2,2)
        x=func.relu(self.conv3(x))
        x=func.max_pool2d(x,2,2)
        x=x.view(-1,4*4*64)
        x=func.relu(self.fully1(x))
        x=self.dropout1(x)
        x=self.fully2(x)
        return x

model=LeNet().to(device)
criterion=nn.CrossEntropyLoss()
optimizer=torch.optim.Adam(model.parameters(),lr=0.001)
```



Step 5:

```
epochs=12
loss_history=[]
correct_history=[]
val_loss_history=[]
val_correct_history=[]
for e in range(epochs):
    loss=0.0
    correct=0.0
    val_loss=0.0
    val_correct=0.0
    for input,labels in training_loader:
        input=input.to(device)
        labels=labels.to(device)
        outputs=model(input)
        loss1=criterion(outputs,labels)
        optimizer.zero_grad()
        loss1.backward()
        optimizer.step()
        _,preds=torch.max(outputs,1)
        loss+=loss1.item()
        correct+=torch.sum(preds==labels.data)
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked in a slightly offset manner, creating a sense of balance and harmony. The background is a soft, out-of-focus landscape with a light sky and a calm body of water that reflects the stones and the sky.

```
else:
    with torch.no_grad():
        for val_input, val_labels in validation_loader:

            val_input=val_input.to(device)
            val_labels=val_labels.to(device)
            val_outputs=model(val_input)
            val_loss1=criterion(val_outputs, val_labels)

            _, val_preds=torch.max(val_outputs, 1)
            val_loss+=val_loss1.item()
            val_correct+=torch.sum(val_preds==val_labels.data)

        epoch_loss=loss/len(training_loader)
        epoch_acc=correct.float()/len(training_loader)
        loss_history.append(epoch_loss)
        correct_history.append(epoch_acc)
        val_epoch_loss=val_loss/len(validation_loader)
        val_epoch_acc=val_correct.float()/len(validation_loader)



        val_loss_history.append(val_epoch_loss)
        val_correct_history.append(val_epoch_acc)
        print('training_loss:{:.4f},{:.4f}'.format(epoch_loss, epoch_acc.item()))
        print('validation_loss:{:.4f},{:.4f}'.format(val_epoch_loss, val_epoch_acc.item()))
```

Step 6:

```
plt.plot(loss_history, label='Training Loss')
plt.plot(val_loss_history, label='Validation Loss')
plt.legend()
plt.show()

plt.plot(correct_history, label='Training accuracy')
plt.plot(val_correct_history, label='Validation accuracy')

plt.legend()
plt.show()
```






Step 7:

```
url='https://akm-img-a-  
in.tosshub.com/indiatoday/images/story/201810/white_stork  
.jpeg?B2LINO47jclclb3QCW.Bj9nto934Lox4'  
response=requests.get(url,stream=True)  
img=Image.open(response.raw)  
img=transform1(img)  
image1=img.to(device).unsqueeze(0)  
output=model(image1)  
_,pred=torch.max(output,1)  
print(classes[pred.item()])
```

Step 8:

```
dataiter=iter(validation_loader)
images,labels=dataiter.next()
images_=images.to(device)
labels=labels.to(device)
output=model(images_)
_,preds=torch.max(output,1)
fig=plt.figure(figsize=(25,4))
for idx in np.arange(20):
    ax=fig.add_subplot(2,10,idx+1,xticks=[],yticks=[])
    plt.imshow(im_convert(images[idx]))

    ax.set_title("{} ({} )".format(str(classes[preds[idx].item()]),str(classes[labels[idx].item()]),color=("green" if classes[preds[idx]]==classes[labels[idx]] else "red")))
plt.show()
```





Any Questions?