

“Don't expect to be easy to work, because if you're easy to do, your heart will be flattered and arrogant.”
_ Buddha



DIP: Multiscale Techniques

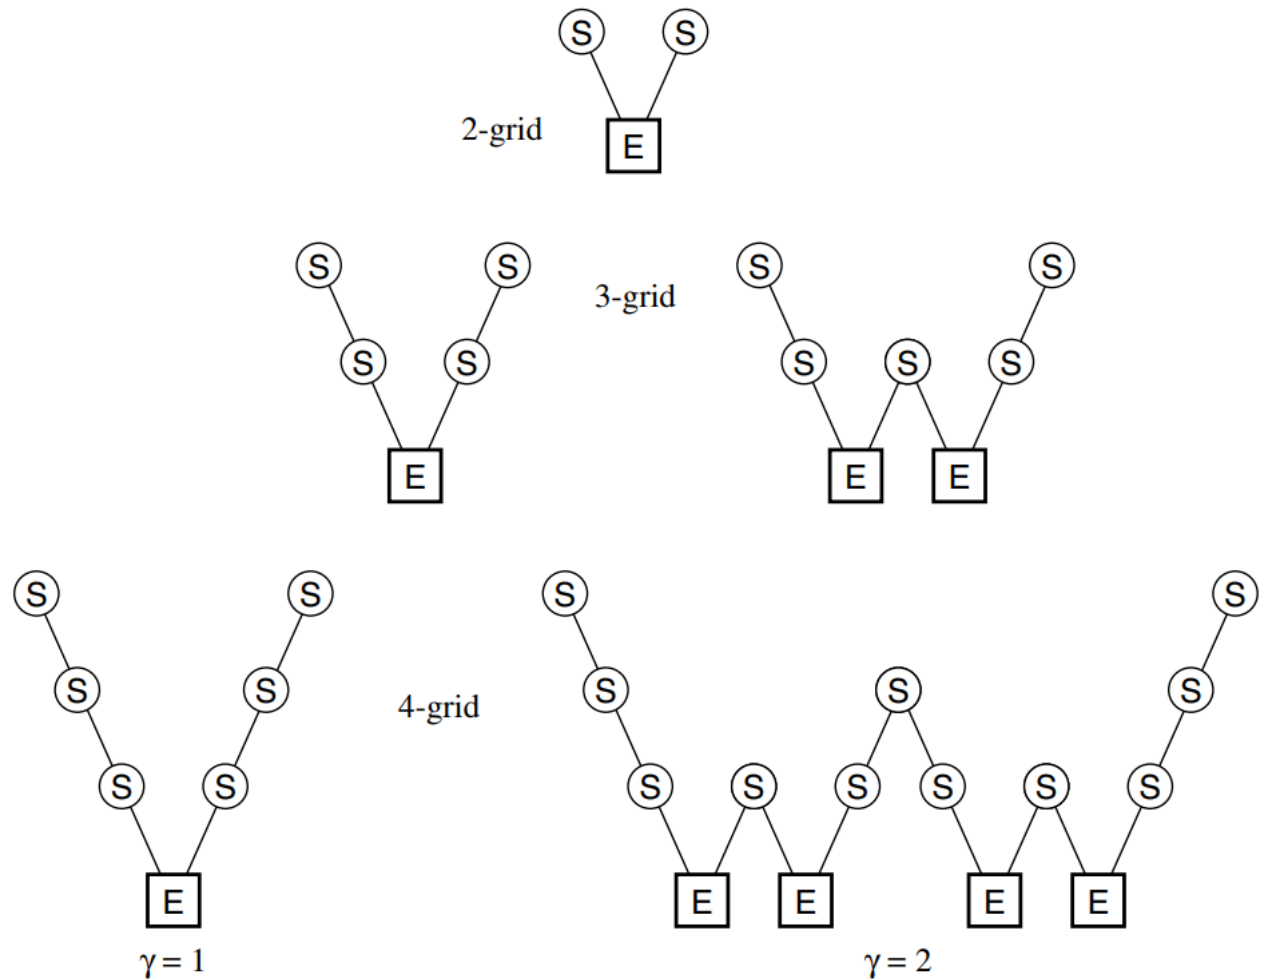
Presenter: Dr. Ha Viet Uyen Synh.



Part A

MULTIGRID TECHNIQUES

Structure of multigrid cycles



S denotes smoothing, while **E** denotes exact solution on the coarsest grid.



Part C

POOLING LAYER



Pooling Layer

Pooling layers section would reduce the number of parameters when the images are too large. Spatial pooling also called subsampling or down sampling which reduces the dimensionality of each map but retains important information.

Spatial pooling can be of different types:


- Max Pooling
- Average Pooling
- Sum Pooling

Max pooling takes the largest element from the rectified feature map. Taking the largest element could also take the average pooling. Sum of all elements in the feature map call as sum pooling.




Example

Max Pooling




Avg Pooling




<https://indoml.com>

Input



Max Pool



<https://indoml.com>



MaxPool2d

Syntax:

```
torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1,  
return_indices=False, ceil_mode=False)
```

The parameters `kernel_size`, `stride`, `padding`, `dilation` can either be:

- a single int – in which case the same value is used for the height and width dimension
- a tuple of two ints – in which case, the first int is used for the height dimension, and the second int for the width dimension

- Input: (N, C, H_{in}, W_{in}) or (C, H_{in}, W_{in})
- Output: (N, C, H_{out}, W_{out}) or (C, H_{out}, W_{out}) , where

$$H_{out} = \left\lfloor \frac{H_{in} + 2 * \text{padding}[0] - \text{dilation}[0] \times (\text{kernel_size}[0] - 1) - 1}{\text{stride}[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 * \text{padding}[1] - \text{dilation}[1] \times (\text{kernel_size}[1] - 1) - 1}{\text{stride}[1]} + 1 \right\rfloor$$

Example

```
# Import the required libraries
import torch
import torch.nn as nn
'''input of size = [N,C,H, W] or [C,H, W]
N==>batch size,
C==> number of channels,
H==> height of input planes in pixels,
W==> width in pixels.
'''


input = torch.empty(3, 4, 4).random_(256)
print("Input Tensor:", input)
print("Input Size:",input.size())
```

```
Input Tensor: tensor([
  [[ 62., 215., 33., 70.],
   [220., 228., 173., 81.],
   [173., 117., 19., 90.],
   [241., 63., 101., 218.]],

  [[221., 85., 104., 69.],
   [144., 63., 149., 187.],
   [132., 252., 152., 211.],
   [244., 5., 55., 191.]],

  [[195., 127., 247., 175.],
   [ 56., 61., 105., 72.],
   [ 52., 233., 20., 147.],
   [ 23., 184., 2., 114.] ]])
```

```
Input Size: torch.Size([3, 4, 4])
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked in a slightly offset manner, creating a sense of balance and harmony. The background is a soft, out-of-focus landscape with a light sky and a calm body of water.

```
# pool of square window of size=3, stride=1
pooling1 = nn.MaxPool2d(3, stride=1)
```


```
# Perform Max Pool
output = pooling1(input)
print("Output Tensor:", output)
print("Output Size:", output.size())
```

```
Output Tensor: tensor([
  [[228., 228.],
   [241., 228.]],
```

```
  [[252., 252.],
   [252., 252.]],
```

```
  [[247., 247.],
   [233., 233.]])
```

```
Output Size: torch.Size([3, 2, 2])
```



```
# pool of non-square window
pooling2 = nn.MaxPool2d((2, 1), stride=(1, 2))
print("Output Kernel:",pooling2)
# Perform Max Pool
output = pooling2(input)
print("Output Tensor:", output)
print("Output Size:",output.size())
```

```
Output Kernel: MaxPool2d(kernel_size=(2, 1), stride=(1,
2), padding=0, dilation=1, ceil_mode=False)
```

```
Output Tensor: tensor([
  [[220., 173.],
   [220., 173.],
   [241., 101.]],

  [[221., 149.],
   [144., 152.],
   [244., 152.]],

  [[195., 247.],
   [ 56., 105.],
   [ 52., 20.]]])
```

```
Output Size: torch.Size([3, 3, 2])
```

Example

```
# Import the required libraries
import torch
import torchvision
from PIL import Image
import torchvision.transforms as T
import torch.nn.functional as F

# read the input image
img = Image.open('elephant.jpg')

# convert the image to torch tensor
img = T.ToTensor()(img)
print("Original size of Image:", img.size()) #Size([3, 466, 700])

# unsqueeze to make 4D
img = img.unsqueeze(0)

# define max pool with square window of size=4, stride=1
pool = torch.nn.MaxPool2d(4, 1)
img = pool(img)
img = img.squeeze(0)
print("Size after MaxPool:",img.size())
img = T.ToPILImage()(img)
img.show()
```





AvgPool2d

Syntax:

```
torch.nn.AvgPool2d(kernel_size, stride=None, padding=0,  
ceil_mode=False, count_include_pad=True, divisor_override=None)
```

The parameters `kernel_size`, `stride`, `padding` can either be:

- a single int – in which case the same value is used for the height and width dimension
 - a tuple of two ints – in which case, the first int is used for the height dimension, and the second int for the width dimension
- Input: (N, C, H_{in}, W_{in}) or (C, H_{in}, W_{in}) .
 - Output: (N, C, H_{out}, W_{out}) or (C, H_{out}, W_{out}) , where

$$H_{out} = \left\lfloor \frac{H_{in} + 2 \times \text{padding}[0] - \text{kernel_size}[0]}{\text{stride}[0]} + 1 \right\rfloor$$

$$W_{out} = \left\lfloor \frac{W_{in} + 2 \times \text{padding}[1] - \text{kernel_size}[1]}{\text{stride}[1]} + 1 \right\rfloor$$

Example

```
# Import the required libraries
import torch
import torchvision
from PIL import Image
import torchvision.transforms as T
import torch.nn.functional as F

# read the input image
img = Image.open('panda.jpg')

# convert the image to torch tensor
img = T.ToTensor()(img)
print("Original size of Image:", img.size())#Size([3, 466, 700])

# unsqueeze to make 4D
img = img.unsqueeze(0)

# define avg pool with square window of size=4, stride=1
pool = torch.nn.AvgPool2d(4, 1)
img = pool(img)
img = img.squeeze(0)
print("Size after AvgPool:",img.size())
img = T.ToPILImage()(img)
img.show()
```






Part D

APPLICATIONS _ U-NET VS W-NET

A vertical stack of five smooth, dark, rounded stones on a reflective surface, likely water. The stones are stacked in a slightly offset manner, and their reflection is visible in the calm water below. The background is a soft, light blue gradient.


What is U-Net?

The U-NET was developed by Olaf Ronneberger et al. for Bio Medical Image Segmentation.


The architecture contains two paths.

- The first path is the *contraction path* (also called as the encoder) which is used to capture the context in the image. The encoder is just a traditional stack of convolutional and max pooling layers.
- The second path is the *symmetric expanding path* (also called as the decoder) which is used to enable precise localization using transposed convolutions.

U-Net Segmentation

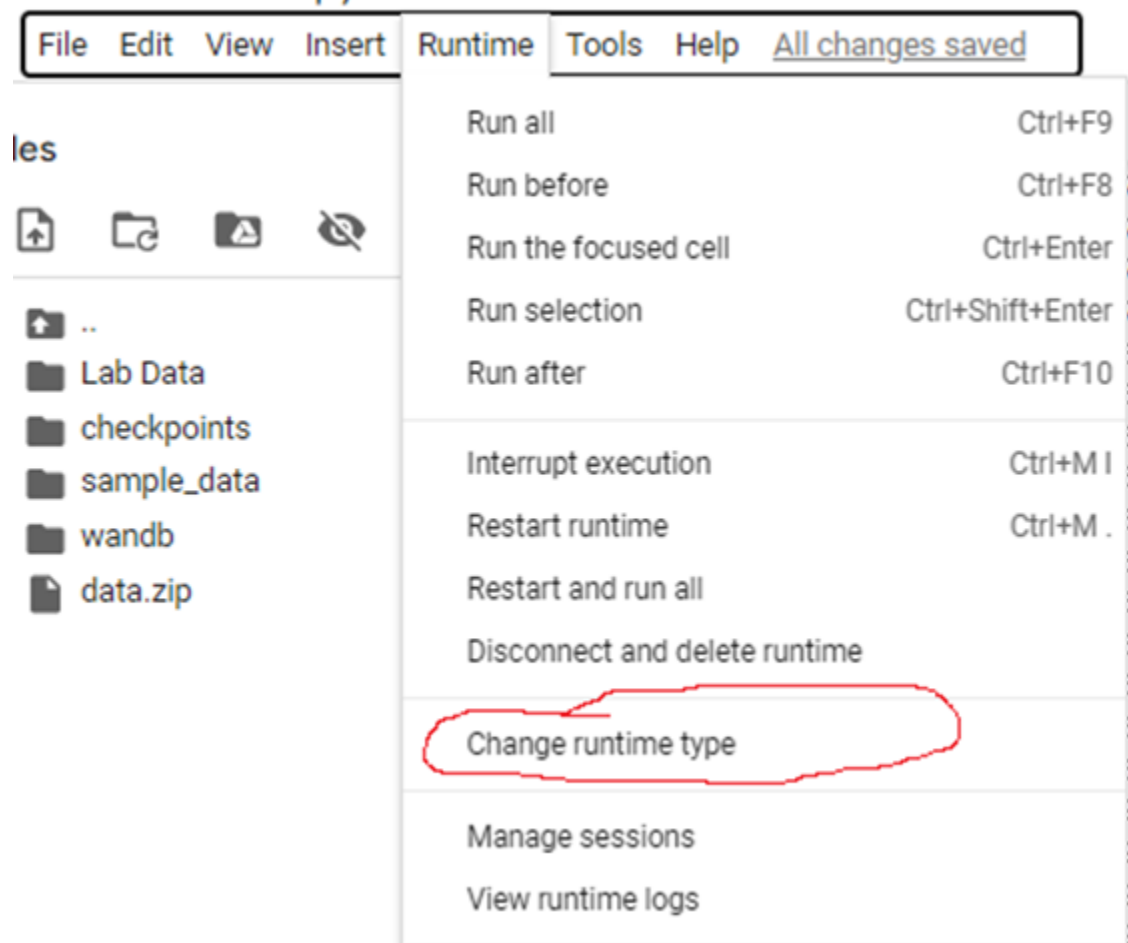


W-Net Segmentation



1. Set up environment for Google Colab

Runtime -> Change runtime type to enable GPU on GG colab notebook



Select GPU of the Hardware accelerator tab and Save notebook

Notebook settings

Hardware accelerator

GPU



Want access to premium GPUs?

[Purchase additional compute units here.](#)

Runtime shape

High-RAM

☐

Omit code cell output when saving this notebook

Cancel

Save

A stack of smooth, dark stones is positioned on the left side of the slide, resting on a reflective surface that shows their reflection. The stones are stacked vertically, with the top stone being the most prominent. The background is a light, neutral color.

2. Training Data - Carvana Image Masking Challenge

We select a subset of 544 images from the original dataset Carvana Image Masking Challenge . You can train the segmentation model with the entire data. However, the training process will take longer time.

This dataset contains a large number of car images (as .jpg files). Each car has exactly 16 images, each one taken at different angles. Each car has a unique id and images are named according to id_01.jpg, id_02.jpg ... id_16.jpg. In addition to the images, you are also provided some basic metadata about the car make, model, year, and trim.

For the training set, you are provided a .gif file that contains the manually cutout mask for each image. The competition task is to automatically segment the cars in the images in the test set folder. To deter hand labeling, we have supplemented the test set with car images that are ignored in scoring.

Note 1: in Lab Data folder

imgs folder: this folder contains the training set images

masks folder: this folder contains the training set masks in .gif format

Note 2: For the other dataset, the format of mask images should be .gif format and binary value (0 and 255) if there are 2 classes



Load datasets

To compare the training process of 2 models: U-Net and W-Net, we use wandb library for visualization. You can create your own account if you want

```
!wget --load-  
cookies /tmp/cookies.txt "https://docs.google.com/uc?expo  
rt=download&confirm=$(wget --quiet --save-  
cookies /tmp/cookies.txt --keep-session-cookies --no-  
check-  
certificate 'https://docs.google.com/uc?export=download&i  
d=1Ju8sx3gpemfY3YILLs_ZNxrLF_05b-1b' -O- | sed -  
rn 's/.*confirm=([0-9A-Za-  
z_]+).*/\1\n/p')&id=1Ju8sx3gpemfY3YILLs_ZNxrLF_05b-1b" -  
O data.zip && rm -rf /tmp/cookies.txt  
!unzip /content/data.zip -d /content  
  
!pip install wandb
```




Import the libraries needed


```
import logging
from os import listdir
from os.path import splitext
from pathlib import Path
import numpy as np
import torch
import wandb
from torch import Tensor
from PIL import Image
from torch.utils.data import Dataset
import torch.nn.functional as F
from tqdm import tqdm
import matplotlib.pyplot as plt
import cv2
import numpy as np
from torch import optim
from torch.utils.data import DataLoader, random_split
from tqdm import tqdm

import torch
import torch.nn as nn
import torch.nn.functional as F
```

Create the BasicDataset class

```
class BasicDataset(Dataset):
    def __init__(self, images_dir: str, masks_dir: str, scale: float = 1.0, mask_suffix: str = ''):
        self.images_dir = Path(images_dir)
        self.masks_dir = Path(masks_dir)
        assert 0 < scale <= 1, 'Scale must be between 0 and 1'
        self.scale = scale
        self.mask_suffix = mask_suffix
        self.ids = [splitext(file)[0] for file in listdir(images_dir) if not file.startswith('.')]
        if not self.ids:
            raise RuntimeError(f'No input file found in {images_dir}, make sure you put your images there')
        logging.info(f'Creating dataset with {len(self.ids)} examples')

    def __len__(self):
        return len(self.ids)
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked horizontally, with the top stone being the most prominent. The background is a soft, out-of-focus blue and white, suggesting a sky or water surface. The overall aesthetic is calm and minimalist.


```
@staticmethod
def preprocess(pil_img, scale, is_mask):
    w, h = pil_img.size
    newW, newH = int(scale * w), int(scale * h)
    assert newW > 0 and newH > 0, 'Scale is too small, resized images would have no pixel'
    pil_img = pil_img.resize((newW, newH), resample=Image.NEAREST if is_mask else Image.BICUBIC)
    img_ndarray = np.asarray(pil_img)

    if not is_mask:
        if img_ndarray.ndim == 2:
            img_ndarray = img_ndarray[np.newaxis, ...]
        else:
            img_ndarray = img_ndarray.transpose((2, 0, 1))

    img_ndarray = img_ndarray / 255

    return img_ndarray

@staticmethod
def load(filename):
    ext = splitext(filename)[1]
    if ext == '.npy':
        return Image.fromarray(np.load(filename))
    elif ext in ['.pt', '.pth']:
        return Image.fromarray(torch.load(filename).numpy())
    else:
        return Image.open(filename)
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked horizontally, and their reflection is visible in the water below. The background is a soft, out-of-focus blue and white, suggesting a sky or water surface.

```
def __getitem__(self, idx):
    name = self.ids[idx]
    mask_file = list(self.masks_dir.glob(name + self.mask_suffix +
        '.*'))
    img_file = list(self.images_dir.glob(name + '.*'))

    assert len(img_file) == 1, f'Either no image or multiple images found for the ID {name}: {img_file}'
    assert len(mask_file) == 1, f'Either no mask or multiple masks found for the ID {name}: {mask_file}'
    mask = self.load(mask_file[0])
    img = self.load(img_file[0])
    assert img.size == mask.size, \
        f'Image and mask {name} should be the same size, but are {img.size} and {mask.size}'
    img = self.preprocess(img, self.scale, is_mask=False)
    mask = self.preprocess(mask, self.scale, is_mask=True)
    return {
        'image': torch.as_tensor(img.copy()).float().contiguous(),
        'mask': torch.as_tensor(mask.copy()).long().contiguous()
    }

class CarvanaDataset(BasicDataset):
    def __init__(self, images_dir, masks_dir, scale=1):
        super().__init__(images_dir, masks_dir, scale, mask_suffix='_mask')
```


Define Dice coefficient for training model

```
def dice_coeff(input: Tensor, target: Tensor, reduce_batch_first: bool = False, epsilon=1e-6):
    # Average of Dice coefficient for all batches, or for a single mask
    assert input.size() == target.size()
    if input.dim() == 2 and reduce_batch_first:
        raise ValueError(f'Dice: asked to reduce batch but got tensor without batch dimension (shape {input.shape})')

    if input.dim() == 2 or reduce_batch_first:
        inter = torch.dot(input.reshape(-1), target.reshape(-1))

        sets_sum = torch.sum(input) + torch.sum(target)
        if sets_sum.item() == 0:
            sets_sum = 2 * inter

        return (2 * inter + epsilon) / (sets_sum + epsilon)
    else:
        # compute and average metric for each batch element
        dice = 0
        for i in range(input.shape[0]):
            dice += dice_coeff(input[i, ...], target[i, ...])
        return dice / input.shape[0]
```

A stack of smooth, dark stones is positioned on the left side of the image, resting on a highly reflective surface that creates a clear mirror image of the stones. The background is a soft, out-of-focus light blue and white, suggesting a calm body of water under a bright sky. The stones are stacked horizontally, with each stone slightly offset from the one below it, creating a sense of depth and balance.

```
def multiclass_dice_coeff(input: Tensor, target: Tensor,
                           reduce_batch_first: bool = False, epsilon=1e-6):
    # Average of Dice coefficient for all classes
    assert input.size() == target.size()
    dice = 0
    for channel in range(input.shape[1]):
        dice += dice_coeff(input[:, channel, ...], target
                          [:, channel, ...], reduce_batch_first, epsilon)

    return dice / input.shape[1]


def dice_loss(input: Tensor, target: Tensor, multiclass:
              bool = False):
    # Dice loss (objective to minimize) between 0 and 1
    assert input.size() == target.size()
    fn = multiclass_dice_coeff if multiclass else dice_coeff
    return 1 - fn(input, target, reduce_batch_first=True)
```

3. Build the main module

```
class DoubleConv(nn.Module):
    """(convolution => [BN] => ReLU) * 2"""

    def __init__(self, in_channels, out_channels, mid_channels=None):
        super().__init__()
        if not mid_channels:
            mid_channels = out_channels
        self.double_conv = nn.Sequential(
            nn.Conv2d(in_channels, mid_channels, kernel_size=3, padding=1, bias=False),
            nn.BatchNorm2d(mid_channels),
            nn.ReLU(inplace=True),
            nn.Conv2d(mid_channels, out_channels, kernel_size=3, padding=1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True)
        )

    def forward(self, x):
        return self.double_conv(x)
```




```
class Down(nn.Module):
    """Downscaling with maxpool then double conv"""

    def __init__(self, in_channels, out_channels):
        super().__init__()
        self.maxpool_conv = nn.Sequential(
            nn.MaxPool2d(2),
            DoubleConv(in_channels, out_channels)
        )

    def forward(self, x):
        return self.maxpool_conv(x)


class OutConv(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(OutConv, self).__init__()
        self.conv = nn.Conv2d(in_channels, out_channels,
                                kernel_size=1)

    def forward(self, x):
        return self.conv(x)
```


```
class Up(nn.Module):
    """Upscaling then double conv"""
    def __init__(self, in_channels, out_channels, bilinear=True):
        super().__init__()

        # if bilinear, use the normal convolutions to reduce the number
        # of channels
        if bilinear:
            self.up = nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
            self.conv = DoubleConv(in_channels, out_channels, in_channels // 2)
        else:
            self.up = nn.ConvTranspose2d(in_channels, in_channels // 2, kernel_size=2, stride=2)
            self.conv = DoubleConv(in_channels, out_channels)

    def forward(self, x1, x2):
        x1 = self.up(x1)
        # input is CHW
        diffY = x2.size()[2] - x1.size()[2]
        diffX = x2.size()[3] - x1.size()[3]

        x1 = F.pad(x1, [diffX // 2, diffX - diffX // 2,
                        diffY // 2, diffY - diffY // 2])

        # if you have padding issues, see
        # https://github.com/HaiyongJiang/U-Net-Pytorch-Unstructured-Buggy/commit/0e854509c2cea854e247a9c615f175f76fbb2e3a
        # https://github.com/xiaopeng-liao/Pytorch-UNet/commit/8ebac70e633bac59fc22bb5195e513d5832fb3bd
        x = torch.cat([x2, x1], dim=1)
        return self.conv(x)
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked horizontally, with the top stone being the most prominent. The background is a soft, out-of-focus landscape with a light sky and a body of water reflecting the stones.

```
def evaluate(net, dataloader, device):
    net.eval()
    num_val_batches = len(dataloader)
    dice_score = 0

    # iterate over the validation set
    for batch in tqdm(dataloader, total=num_val_batches, desc='Validation round', unit='batch', leave=False):
        image, mask_true = batch['image'], batch['mask']
        # move images and labels to correct device and type
        image = image.to(device=device, dtype=torch.float32)
        mask_true = mask_true.to(device=device, dtype=torch.long)
        mask_true = F.one_hot(mask_true, net.n_classes).permute(0, 3, 1, 2).float()

        with torch.no_grad():
            # predict the mask
            mask_pred = net(image)

            # convert to one-hot format
            if net.n_classes == 1:
                mask_pred = (F.sigmoid(mask_pred) > 0.5).float()
                # compute the Dice score
                dice_score += dice_coeff(mask_pred, mask_true, reduce_batch_first=False)
            else:
                mask_pred = F.one_hot(mask_pred.argmax(dim=1), net.n_classes).permute(0, 3, 1, 2).float()
                # compute the Dice score, ignoring background
                dice_score += multiclass_dice_coeff(mask_pred[:, 1:, ...], mask_true[:, 1:, ...], reduce_batch_first=False)

    net.train()

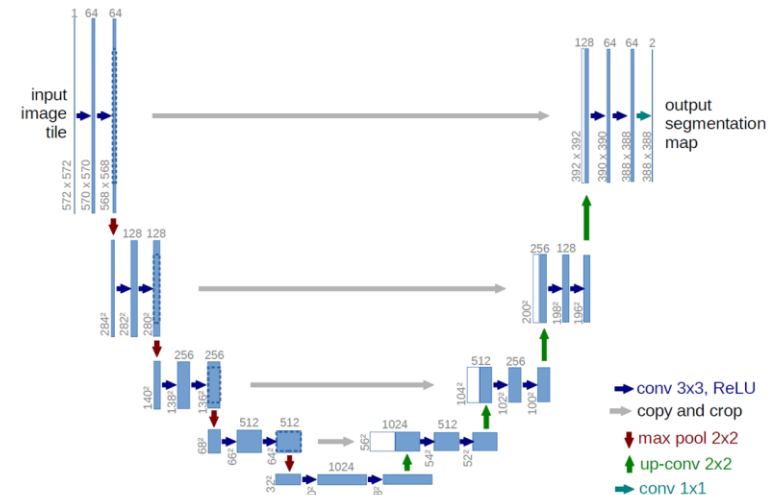
    # Fixes a potential division by zero error
    if num_val_batches == 0:
        return dice_score
    return dice_score / num_val_batches
```

4. Create the U-Net class

```
class UNet(nn.Module):
    def __init__(self, n_channels, n_classes, bilinear=False):
        super(UNet, self).__init__()
        self.n_channels = n_channels
        self.n_classes = n_classes
        self.bilinear = bilinear

        self.inc = DoubleConv(n_channels, 64)
        self.down1 = Down(64, 128)
        self.down2 = Down(128, 256)
        self.down3 = Down(256, 512)
        factor = 2 if bilinear else 1
        self.down4 = Down(512, 1024 // factor)
        self.up1 = Up(1024, 512 // factor, bilinear)
        self.up2 = Up(512, 256 // factor, bilinear)
        self.up3 = Up(256, 128 // factor, bilinear)
        self.up4 = Up(128, 64, bilinear)
        self.outc = OutConv(64, n_classes)


    def forward(self, x):
        x1 = self.inc(x)
        x2 = self.down1(x1)
        x3 = self.down2(x2)
        x4 = self.down3(x3)
        x5 = self.down4(x4)
        x = self.up1(x5, x4)
        x = self.up2(x, x3)
        x = self.up3(x, x2)
        x = self.up4(x, x1)
        logits = self.outc(x)
        return logits
```



5. Build the Train_Net function

```
def train_net(net,
              device,
              epochs: int = 5,
              batch_size: int = 1,
              learning_rate: float = 1e-5,
              val_percent: float = 0.1,
              save_checkpoint: bool = True,
              img_scale: float = 0.5,
              amp: bool = False):
    # 1. Create dataset
    try:
        dataset = CarvanaDataset(dir_img, dir_mask, img_scale)
    except (AssertionError, RuntimeError):
        dataset = BasicDataset(dir_img, dir_mask, img_scale)

    # 2. Split into train / validation partitions
    n_val = int(len(dataset) * val_percent)
    n_train = len(dataset) - n_val
    train_set, val_set = random_split(dataset, [n_train, n_val],
                                       generator=torch.Generator().manual_seed(0))
```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked horizontally, and their reflection is visible in the water below. The background is a soft, out-of-focus blue and white gradient.

```
# 3. Create data loaders
loader_args = dict(batch_size=batch_size, num_workers=4, pin_memory=True)
train_loader = DataLoader(train_set, shuffle=True, **loader_args)
val_loader = DataLoader(val_set, shuffle=False, drop_last=True, **loader_args)

# (Initialize logging)
experiment = wandb.init(project='Unet-Lab', anonymous='must')
experiment.config.update(dict(epochs=epochs, batch_size=batch_size,
learning_rate=learning_rate,
val_percent=val_percent, save_checkpoint=save_checkpoint, img_scale=img_scale, amp=amp))

logging.info(f'''Starting training:
    Epochs:          {epochs}
    Batch size:      {batch_size}
    Learning rate:   {learning_rate}
    Training size:   {n_train}
    Validation size: {n_val}
    Checkpoints:     {save_checkpoint}
    Device:          {device.type}
    Images scaling:  {img_scale}
    Mixed Precision: {amp}
''')
```




```
# 4. Set up the optimizer, the loss, the learning rate scheduler
and the loss scaling for AMP
optimizer = optim.RMSprop(net.parameters(), lr=learning_rate, wei
ght_decay=1e-8, momentum=0.9)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, 'max'
, patience=2) # goal: maximize Dice score
grad_scaler = torch.cuda.amp.GradScaler(enabled=amp)
criterion = nn.CrossEntropyLoss()
global_step = 0

# 5. Begin training
for epoch in range(1, epochs+1):
    net.train()
    epoch_loss = 0
    with tqdm(total=n_train, desc=f'Epoch {epoch}/{epochs}', unit='img'
) as pbar:
        for batch in train_loader:
            images = batch['image']
            true_masks = batch['mask']

            assert images.shape[1] == net.n_channels, \
                f'Network has been defined with {net.n_channels} input
channels, but loaded images have {images.shape[1]} channels. Please check t
hat the images are loaded correctly.'

            images = images.to(device=device, dtype=torch.float32)
            true_masks = true_masks.to(device=device, dtype=torch.long)
```

A vertical stack of five smooth, dark, rounded stones sits on a calm, reflective surface. The stones are stacked in a slightly offset manner, creating a sense of balance and harmony. The surface of the water or liquid they sit on is still, creating a clear reflection of the stones and the sky above. The background is a soft, hazy blue, suggesting a clear sky or a distant horizon. The overall mood is peaceful and serene.

```
with torch.cuda.amp.autocast(enabled=amp):
    masks_pred = net(images)
    loss = criterion(masks_pred, true_masks) \
        + dice_loss(F.softmax(masks_pred, dim=
1).float(), F.one_hot(true_masks, net.n_classes).permute(0, 3, 1,
2).float(), multiclass=True)

    optimizer.zero_grad(set_to_none=True)
    grad_scaler.scale(loss).backward()
    grad_scaler.step(optimizer)
    grad_scaler.update()

    pbar.update(images.shape[0])
    global_step += 1
    epoch_loss += loss.item()
    experiment.log({
        'train loss': loss.item(),
        'step': global_step,
        'epoch': epoch
    })
    pbar.set_postfix(**{'loss (batch)': loss.item()})
```



```
# Evaluation round
division_step = (n_train // (10 * batch_size))
if division_step > 0:
    if global_step % division_step == 0:
        histograms = {}
        for tag, value in net.named_parameters():
            tag = tag.replace('/', '.')
            if not torch.isinf(value).any():
                histograms['Weights/' + tag] = wandb.Histogram(value.data.cpu())
            if not torch.isinf(value.grad).any():
                histograms['Gradients/' + tag] = wandb.Histogram(value.grad.data.cpu())

        val_score = evaluate(net, val_loader, device)
        scheduler.step(val_score)

        logging.info('Validation Dice score: {}'.format(
            val_score))

        experiment.log({
            'learning rate': optimizer.param_groups[0]['lr'],
            'validation Dice': val_score,
            'images': wandb.Image(images[0].cpu()),
            'masks': {'true': wandb.Image(true_masks[0].float().cpu()),
                       'pred': wandb.Image(masks_pred.argmax(dim=1)[0].float().cpu())},
            'step': global_step,
            'epoch': epoch,
            **histograms})
```




```
if save_checkpoint:
    Path(dir_checkpoint).mkdir(parents=True, exist_ok=True)
    torch.save(net.state_dict(), str(dir_checkpoint / 'check
point_epoch{}.pth'.format(epoch)))
    logging.info(f'Checkpoint {epoch} saved!')
```

6. Parameter for Training U-Net Model

```
dir_checkpoint = Path('./checkpoints_Unet/')
dir_img = Path('./Lab Data/imgs/')
dir_mask = Path('./Lab Data/masks/')

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
epochs = 5
batch_size = 1
lr = 1e-5
scale = 0.5
val = 10
amp = False
classes = 2
bilinear = False

logging.basicConfig(level=logging.INFO, format='%(levelname)s: %(message)s')
logging.info(f'Using device {device}')

net = UNet(n_channels=3, n_classes= classes, bilinear= bilinear)

logging.info(f'Network:\n'
             f'\t{net.n_channels} input channels\n'
             f'\t{net.n_classes} output channels (classes)\n'
             f'\t{"Bilinear" if net.bilinear else "Transposed conv"} upscaling')

net.to(device=device)
```



```
try:
    train_net(net=net,
              epochs= epochs,
              batch_size= batch_size,
              learning_rate= lr,
              device=device,
              img_scale= scale,
              val_percent= val / 100,
              amp= amp)
except KeyboardInterrupt:
    torch.save(net.state_dict(), 'INTERRUPTED.pth')
    logging.info('Saved interrupt')
    raise
```

7. Create the W_Net class

```
class WNet(nn.Module):
```

```
def __init__(self, n_channels, n_classes, bilinear=False):
    super(WNet, self).__init__()
    self.n_channels = n_channels
    self.n_classes = n_classes
    self.bilinear = bilinear
```

```
self.inc = DoubleConv(n_channels, 64)
self.down1 = Down(64, 128)
self.down2 = Down(128, 256)
self.down3 = Down(256, 512)
factor = 2 if bilinear else 1
self.down4 = Down(512, 1024 // factor)
self.up1 = Up(1024, 512 // factor, bilinear)
self.up2 = Up(512, 256 // factor, bilinear)
self.up3 = Up(256, 128 // factor, bilinear)
self.up4 = Up(128, 64, bilinear)
self.outc = OutConv(64, n_classes)
```


```
# Wnet #####
```

```
def forward(self, x):
    x1 = self.inc(x)
    x2 = self.down1(x1)
    x3 = self.down2(x2)
    x4 = self.down3(x3)
    x5 = self.down4(x4)
```

```
x_up1 = self.up1(x5, x4)
x_up2 = self.up2(x_up1, x3)
```

```
x4_4 = self.down3(x_up2)
x5_5 = self.down4(x4_4)
```

```
x_up11 = self.up1(x5_5, x4_4)
x_up22 = self.up2(x_up11, x_up2)
x_up33 = self.up3(x_up22, x2)
x_up44 = self.up4(x_up33, x1)
logits = self.outc(x_up44)
return logits
```





8. Parameter for training W_Net Model

```
dir_checkpoint = Path('./checkpoints_Wnet/')
```

```
dir_img = Path('./Lab Data/imgs/')
```

```
dir_mask = Path('./Lab Data/masks/')
```

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
epochs = 5
```

```
batch_size = 1
```

```
lr = 1e-5
```

```
scale = 0.5
```

```
val = 10
```

```
amp = False
```


```
classes = 2
```

```
bilinear = False
```

9. Build the Train_WNet function

```
def train_Wnet(net,
               device,
               epochs: int = 5,
               batch_size: int = 1,
               learning_rate: float = 1e-5,
               val_percent: float = 0.1,
               save_checkpoint: bool = True,
               img_scale: float = 0.5,
               amp: bool = False):
    # 1. Create dataset
    try:
        dataset = CarvanaDataset(dir_img, dir_mask, img_scale)
    except (AssertionError, RuntimeError):
        dataset = BasicDataset(dir_img, dir_mask, img_scale)

    # 2. Split into train / validation partitions
    n_val = int(len(dataset) * val_percent)
    n_train = len(dataset) - n_val
    train_set, val_set = random_split(dataset, [n_train, n_val],
                                       generator=torch.Generator().manual_seed(0))
```



```
# 3. Create data loaders
loader_args = dict(batch_size=batch_size, num_workers=4, pin_memory=True)
train_loader = DataLoader(train_set, shuffle=True, **loader_args)
val_loader = DataLoader(val_set, shuffle=False, drop_last=True, **loader_args)

# (Initialize logging)
experiment = wandb.init(project='Wnet-Lab', resume='allow', anonymous='must')
experiment.config.update(dict(epochs=epochs, batch_size=batch_size, learning_rate=learning_rate,
                             val_percent=val_percent, save_checkpoint=save_checkpoint, img_scale=img_scale,
                             amp=amp))

logging.info(f'''Starting training:
Epochs:           {epochs}
Batch size:        {batch_size}
Learning rate:     {learning_rate}
Training size:     {n_train}
Validation size:   {n_val}
Checkpoints:       {save_checkpoint}
Device:            {device.type}
Images scaling:    {img_scale}
Mixed Precision:   {amp}
''')
```



```
# 4. Set up the optimizer, the loss, the learning rate scheduler and the loss scaling for AMP
optimizer = optim.RMSprop(net.parameters(), lr=learning_rate,
weight_decay=1e-8, momentum=0.9)
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, '
max', patience=2) # goal: maximize Dice score
grad_scaler = torch.cuda.amp.GradScaler(enabled=amp)
criterion = nn.CrossEntropyLoss()
global_step = 0
```




```
# 5. Begin training
for epoch in range(1, epochs+1):
    net.train()
    epoch_loss = 0
    with tqdm(total=n_train, desc=f'Epoch {epoch}/{epochs}', unit=
'img') as pbar:
        for batch in train_loader:
            images = batch['image']
            true_masks = batch['mask']

            assert images.shape[1] == net.n_channels, \
                f'Network has been defined with {net.n_channels} i
nput channels, but loaded images have {images.shape[1]} channels. Plea
se check that the images are loaded correctly.'
            images = images.to(device=device, dtype=torch.float32)
            true_masks = true_masks.to(device=device, dtype=torch.
long)

            with torch.cuda.amp.autocast(enabled=amp):
                masks_pred = net(images)
                loss = criterion(masks_pred, true_masks) \
                    + dice_loss(F.softmax(masks_pred, dim=1).fl
oat(), F.one_hot(true_masks, net.n_classes).permute(0, 3, 1, 2).float(
), multiclass=True)
```



```
optimizer.zero_grad(set_to_none=True)
grad_scaler.scale(loss).backward()
grad_scaler.step(optimizer)
grad_scaler.update()

pbar.update(images.shape[0])
global_step += 1
epoch_loss += loss.item()
experiment.log({
    'train loss': loss.item(),
    'step': global_step,
    'epoch': epoch
})
pbar.set_postfix(**{'loss (batch)': loss.item(
    )})
```



```

# Evaluation round
division_step = (n_train // (10 * batch_size))
if division_step > 0:
    if global_step % division_step == 0:
        histograms = {}
        for tag, value in net.named_parameters():
            tag = tag.replace('/', '.')
            if not torch.isinf(value).any():
                histograms['Weights/' + tag] = wandb.Histogram
(value.data.cpu())

            if not torch.isinf(value.grad).any():
                histograms['Gradients/' + tag] = wandb.Histogr
am(value.grad.data.cpu())

        val_score = evaluate(net, val_loader, device)
        scheduler.step(val_score)

        logging.info('Validation Dice score: {}'.format(val_sc
ore))

        experiment.log({
            'learning rate': optimizer.param_groups[0]['lr'],
            'validation Dice': val_score,
            'images': wandb.Image(images[0].cpu()),
            'masks': {
                'true': wandb.Image(true_masks[0].float().cpu(
)),
                'pred': wandb.Image(masks_pred.argmax(dim=1)[0
].float().cpu()),
            },
            'step': global_step,
            'epoch': epoch,
            **histograms
        })
    if save_checkpoint:
        Path(dir_checkpoint).mkdir(parents=True, exist_ok=True)
        torch.save(net.state_dict(), str(dir_checkpoint / 'checkpoint_epoc
h{}.pth'.format(epoch)))
        logging.info(f'Checkpoint {epoch} saved!')

```

A stack of smooth, dark stones is shown on the left side of the image, resting on a reflective surface. The stones are stacked horizontally, and their reflection is visible in the water below. The background is a soft, out-of-focus blue and white, suggesting a calm body of water under a clear sky.

```
logging.basicConfig(level=logging.INFO, format='%(levelname)s: %(message)s')
logging.info(f'Using device {device}')


net = WNet(n_channels=3, n_classes= classes, bilinear= bilinear)

logging.info(f'Network:\n'
             f'\t{net.n_channels} input channels\n'
             f'\t{net.n_classes} output channels (classes)\n'
             f'\t{"Bilinear" if net.bilinear else "Transposed conv"} upscaling')


net.to(device=device)

try:
    train_Wnet(net=net,
               epochs= epochs,
               batch_size= batch_size,
               learning_rate= lr,
               device=device,
               img_scale= scale,
               val_percent= val / 100,
               amp= amp)
except KeyboardInterrupt:
    torch.save(net.state_dict(), 'INTERRUPTED.pth')
    logging.info('Saved interrupt')
    raise
```

U-Net




W-Net



10. Predict image

```
!wget --load-  
cookies /tmp/cookies.txt "https://docs.google.com/uc?export=downl  
oad&confirm=$(wget --quiet --save-cookies /tmp/cookies.txt --  
keep-session-cookies --no-check-  
certificate 'https://docs.google.com/uc?export=download&id=1KC90U  
wsY31OXMPq9mtFVRHesVX1zjpAG' -O- | sed -rn 's/.*confirm=([0-9A-  
Za-z_]+).*/\1\n/p')&id=1KC90UwsY31OXMPq9mtFVRHesVX1zjpAG" -  
O data.zip && rm -rf /tmp/cookies.txt  
!unzip /content/data.zip -d /content
```

```
import argparse  
import logging  
import os  
  
import numpy as np  
import torch  
import torch.nn.functional as F  
from PIL import Image  
from torchvision import transforms
```




```
def predict_img(net,
                full_img,
                device,
                scale_factor=1,
                out_threshold=0.5):
    net.eval()
    img = torch.from_numpy(BasicDataset.preprocess(full_img, scale_factor,
    or, is_mask=False))
    img = img.unsqueeze(0)
    img = img.to(device=device, dtype=torch.float32)

    with torch.no_grad():
        output = net(img)
        if net.n_classes > 1:
            probs = F.softmax(output, dim=1)[0]
        else:
            probs = torch.sigmoid(output)[0]

    tf = transforms.Compose([
        transforms.ToPILImage(),
        transforms.Resize((full_img.size[1], full_img.size[0])),
        transforms.ToTensor()
    ])

    full_mask = tf(probs.cpu()).squeeze()

    if net.n_classes == 1:
        return (full_mask > out_threshold).numpy()
    else:
        return F.one_hot(full_mask.argmax(dim=0), net.n_classes).permute(2, 0, 1).numpy()
```

A stack of four smooth, dark, rounded stones is positioned on the left side of the image. They are stacked vertically, with the top stone being the smallest and the bottom one the largest. The stones are resting on a highly reflective, light-colored surface, which creates a clear reflection of the stones below them. The background is a soft, out-of-focus light blue and white, suggesting a sky or a body of water.

```
# checkpoint_dir = "/content/checkpoints_Unet/checkpoint_
epoch5.pth"
checkpoint_dir = "/content/checkpoints/checkpoint_epoch5.
pth"

# image = "/content/predict/6ae670e86620_06.jpg"
image = "/content/predict/6ba36af67cb0_02.jpg"
# image = "/content/predict/6c0cd487abcd_10.jpg"
# image = "/content/predict/6c3470c34408_06.jpg"

# mask = "/content/predict/6ae670e86620_06_mask.gif"
mask = "/content/predict/6ba36af67cb0_02_mask.gif"
# mask = "/content/predict/6c0cd487abcd_10_mask.gif"
# mask = "/content/predict/6c3470c34408_06_mask.gif"

scale = 0.5
```




Visualize the prediction

```
def plot_img_and_mask(img, mask, groundTruth):
    classes = mask.shape[0] if len(mask.shape) > 2 else 1
    fig, ax = plt.subplots(1, classes + 2, figsize=(15, 15))

    ax[0].set_title('Input image')
    ax[0].imshow(img)
    if classes > 1:
        for i in range(classes):
            ax[i + 1].set_title(f'Output mask (class {i + 1})')
            ax[i + 1].imshow(mask[i, :, :])
    else:
        ax[1].set_title(f'Output mask')
        ax[1].imshow(mask)
    ax[3].set_title(f'Groundtruth')
    ax[3].imshow(groundTruth)

    plt.xticks([], plt.yticks([]))
    plt.show()
```

Create and load the trained model

```
## Wnet ###
net = WNet(n_channels=3, n_classes=2, bilinear=False)


## Unet - Uncomment here ###
# net = UNet(n_channels=3, n_classes=2, bilinear=False)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
net.to(device=device)
net.load_state_dict(torch.load(checkpoint_dir, map_location=device))

img = Image.open(image)
groundtruth_mask = Image.open(mask)

mask = predict_img(net=net, full_img=img, scale_factor= scale, out_threshold
= 0.5, device=device)

plot_img_and_mask(img, mask, groundtruth_mask)
```





Any Questions?