



# APT: Image Augmentation

Presenter: Dr. Ha Viet Uyen Synh.



# **INTRODUCTION**

# Warm up \_ Surface Crack dataset

The dataset is generated from **458** high-resolution images (4032x3024 pixel) with the method proposed by Zhang et al (2016).

High-resolution images have variance in terms of surface finish and illumination conditions. No data augmentation in terms of random rotation or flipping is applied.



A vertical stack of five smooth, dark, rounded stones on a reflective surface, likely water. The stones are stacked in a slightly offset manner, and their reflections are visible in the water below. The background is a soft, out-of-focus blue and white.

# Surface Crack dataset

<https://data.mendeley.com/datasets/5y9wdsg2zt/2>

Or <https://www.kaggle.com/datasets/arunrk7/surface-crack-detection>

The dataset is divided into two as negative and positive crack images for image classification. Each class has 20000 images with a total of **40000** images with 227 x 227 pixels with RGB channels.

A vertical stack of five smooth, dark, rounded stones sits on a calm, reflective surface. The stones are stacked in a slightly offset manner, creating a sense of balance and harmony. The surface of the water or liquid they sit on is perfectly still, creating a clear reflection of the stones and the light above. The background is a soft, out-of-focus light blue, suggesting a sky or a distant horizon. The overall mood is peaceful and contemplative.

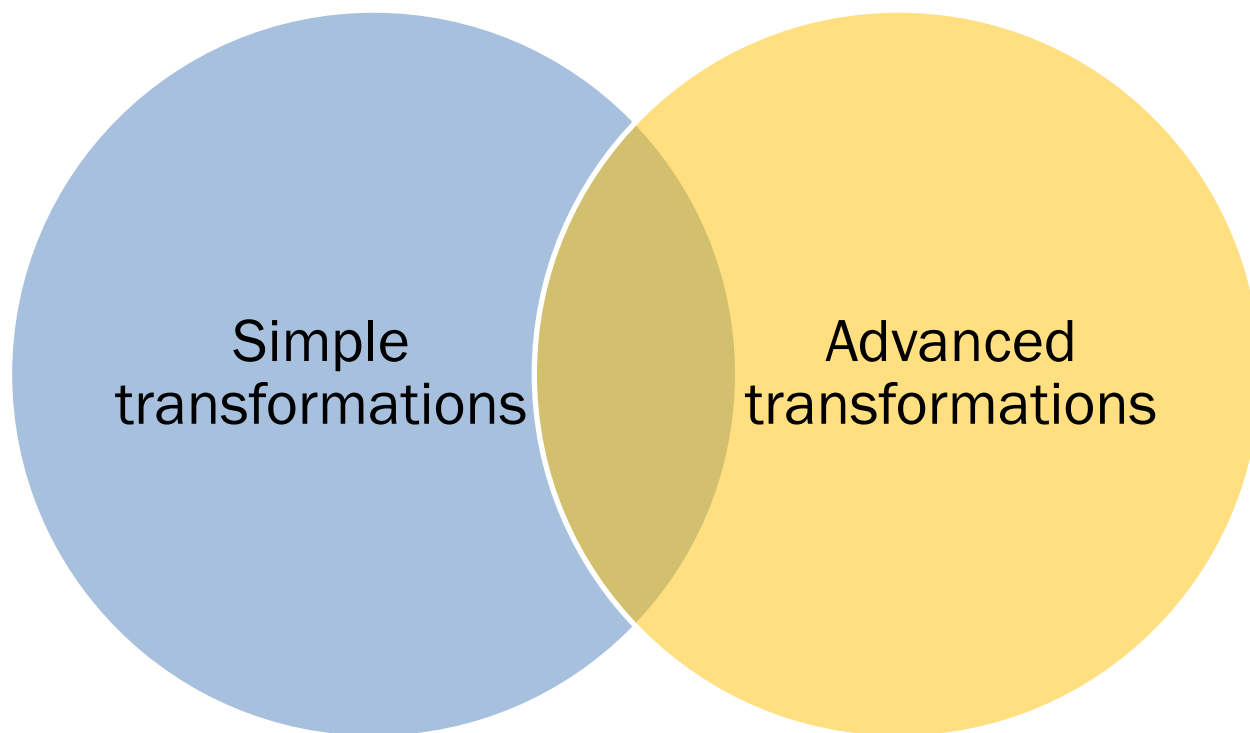
# Image Augmentation

A way to increase the amount of data and make the model more robust

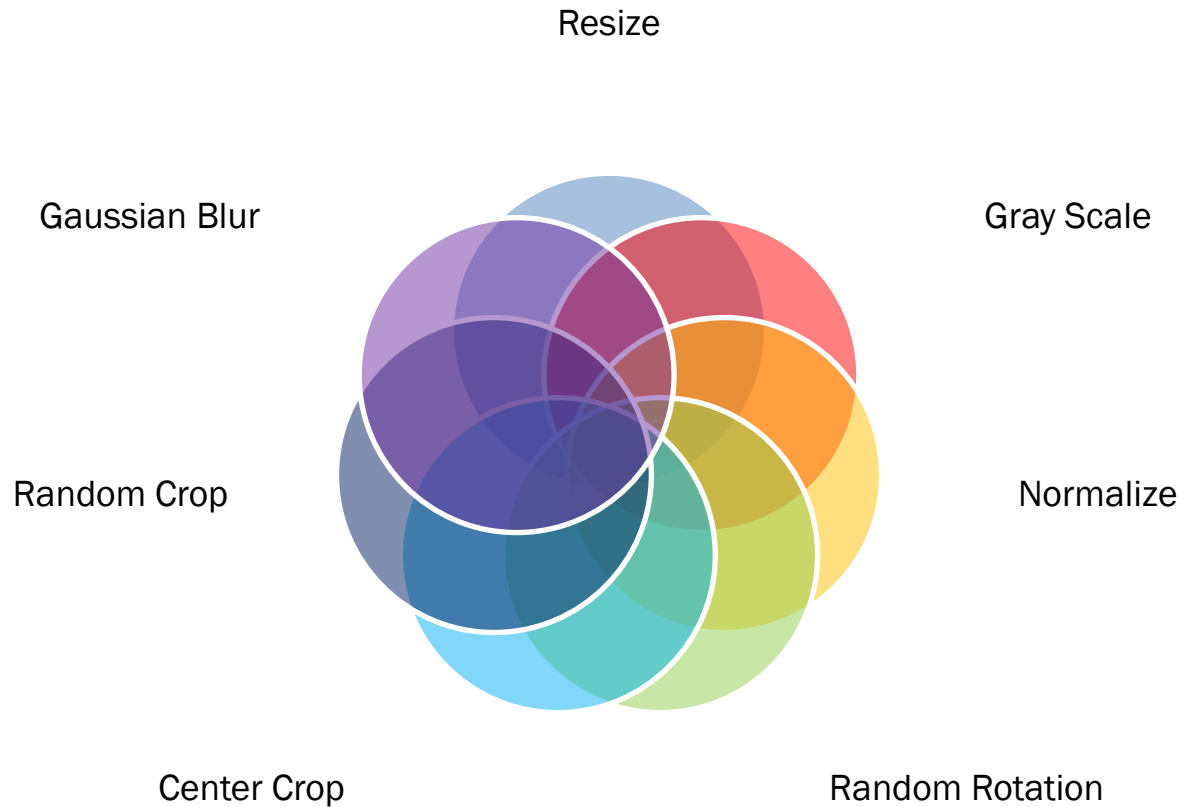
The problem addressed is anomaly detection, which is quite challenging since there is a small volume of data and then, the model is not enough to do all the work alone. The common scenario is to train a two-network model with the normal images available for training and evaluate its performance on the test set, that contains both normal and anomalous images.

The initial hypothesis is that the generative model should capture well the normal distribution but at the same time, it should fail on reconstructing the abnormal samples. How is it possible to verify this hypothesis? We can look at the reconstruction error, which should be higher for abnormal images, while it should be low for the normal samples.

# Image Augmentation techniques



# 1. Simple transformations



A stack of smooth, dark stones is shown on the left side of the slide, resting on a reflective surface. The stones are stacked horizontally, and their reflection is visible in the water below. The background is a soft, out-of-focus blue and white.

# Simple transformations

```
from PIL import Image
from pathlib import Path
import matplotlib.pyplot as plt
import numpy as np
import sys
import torch
import numpy as np
import torchvision.transforms as T

plt.rcParams["savefig.bbox"] = 'tight'
orig_img = Image.open(Path('../input/surface-crack-
detection/Negative/00026.jpg'))
torch.manual_seed(0)
data_path = '../input/surface-crack-detection/'
diz_class = {'Positive': 'Crack', 'Negative': 'No crack'}

np.asarray(orig_img).shape  #(227, 227, 3)
```



# Resize

```
resized_imgs = [T.Resize(size=size)(orig_img) for size in  
[32,128]]  
plot(resized_imgs,col_title=["32x32","128x128"])
```

Original image



32x32



128x128



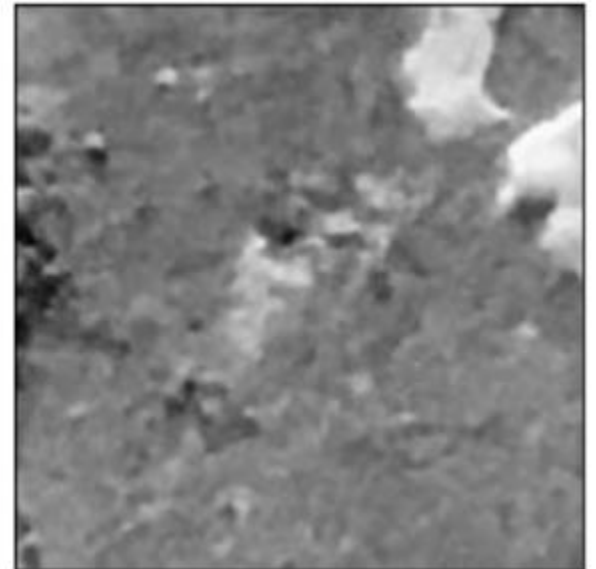
# Gray Scale

```
gray_img = T.Grayscale()(orig_img)  
plot([gray_img], cmap='gray', col_title=["Gray"])
```

Original image



Gray



# Normalize

The normalization can constitute an effective way to speed up the computations in the model based on neural network architecture and learn faster. There are two steps to normalize the images:

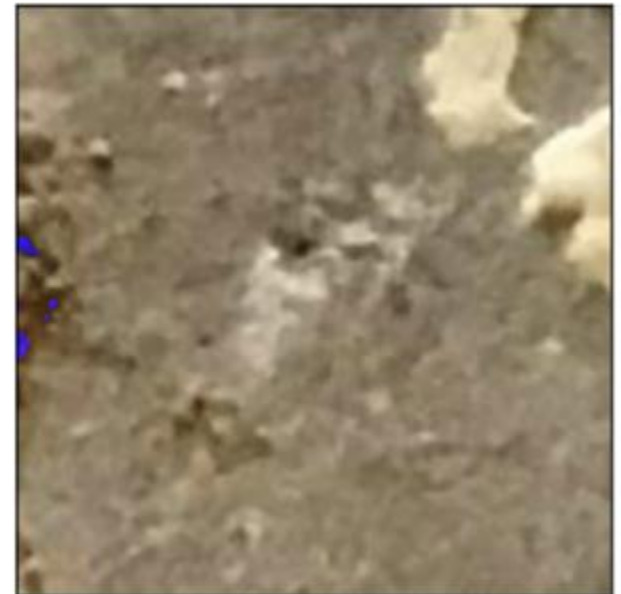
- we subtract the channel mean from each input channel
- we divide it by the channel standard deviation.

```
normalized_img = T.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))(T.ToTensor()(orig_img))  
normalized_img = [T.ToPILImage()(normalized_img)]  
plot(normalized_img, col_title=["Standard normalize"])
```

Original image

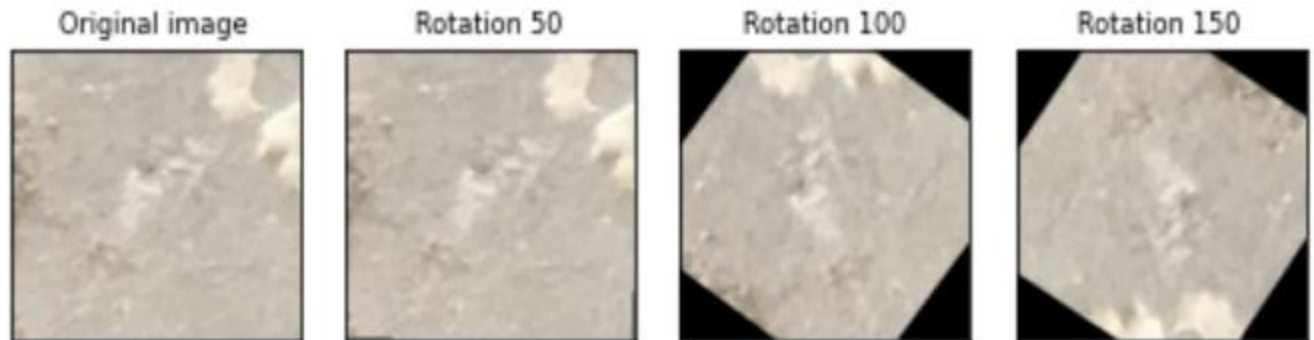


Standard normalize



# Random Rotation

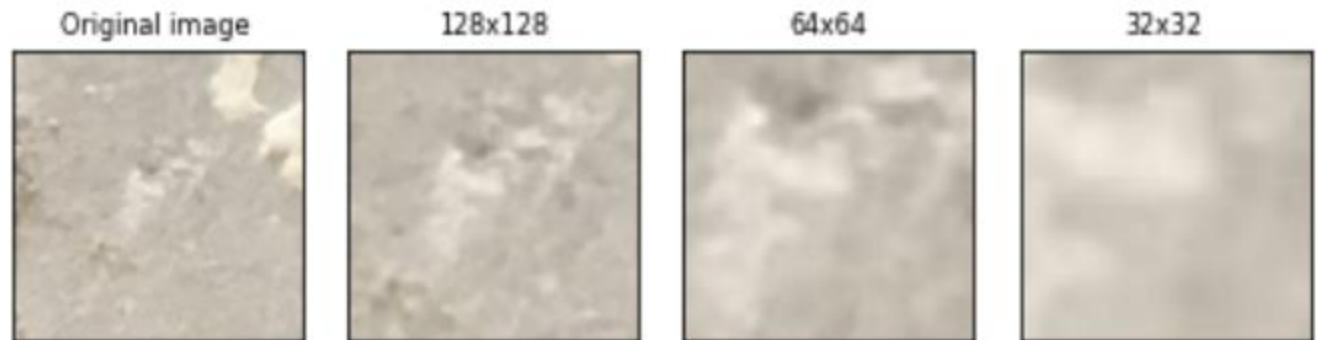
```
rotated_imgs = [T.RandomRotation(degrees=d)(orig_img) for  
d in range(50,151,50)]  
plot(rotated_imgs, col_title=["Rotation 50", "Rotation 100",  
"Rotation 150"])
```



# Center Crop

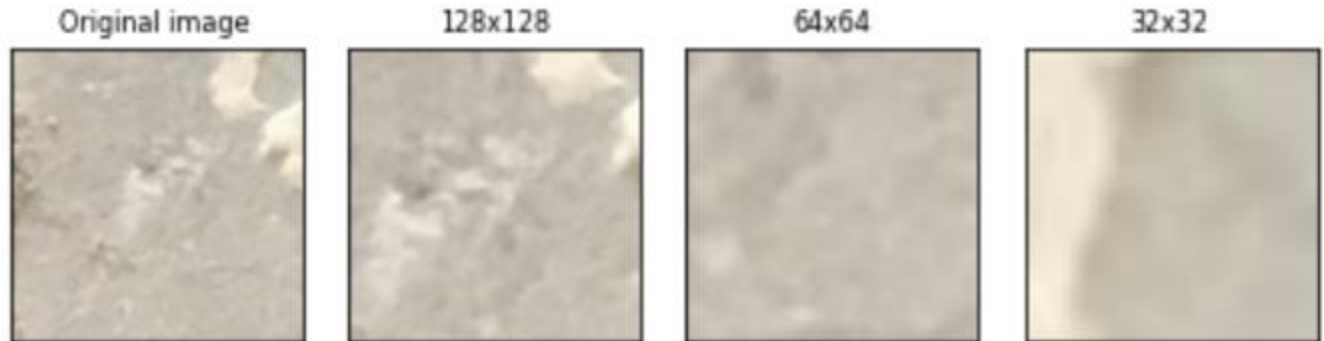
This transformation can be useful when the image has a big background in the borders that isn't necessary at all for the classification task.

```
center_crops = [T.CenterCrop(size=size)(orig_img) for size in (128, 64, 32)]  
plot(center_crops, col_title=['128x128', '64x64', '32x32'])
```



# Random drop

```
random_crops = [T.RandomCrop(size=size)(orig_img) for size in (832, 704, 256)]  
plot(random_crops, col_title=['832x832', '704x704', '256x256'])
```



# Gaussian Blur

This method can be helpful in making the image less clear and distinct and, then, this resulting image is fed into a neural network, which becomes more robust in learning patterns of the samples

```
blurred_imgs = [T.GaussianBlur(kernel_size=(51, 91), sigma=sigma)
                 (orig_img) for sigma in (3,7)]
plot(blurred_imgs)
```

Original image



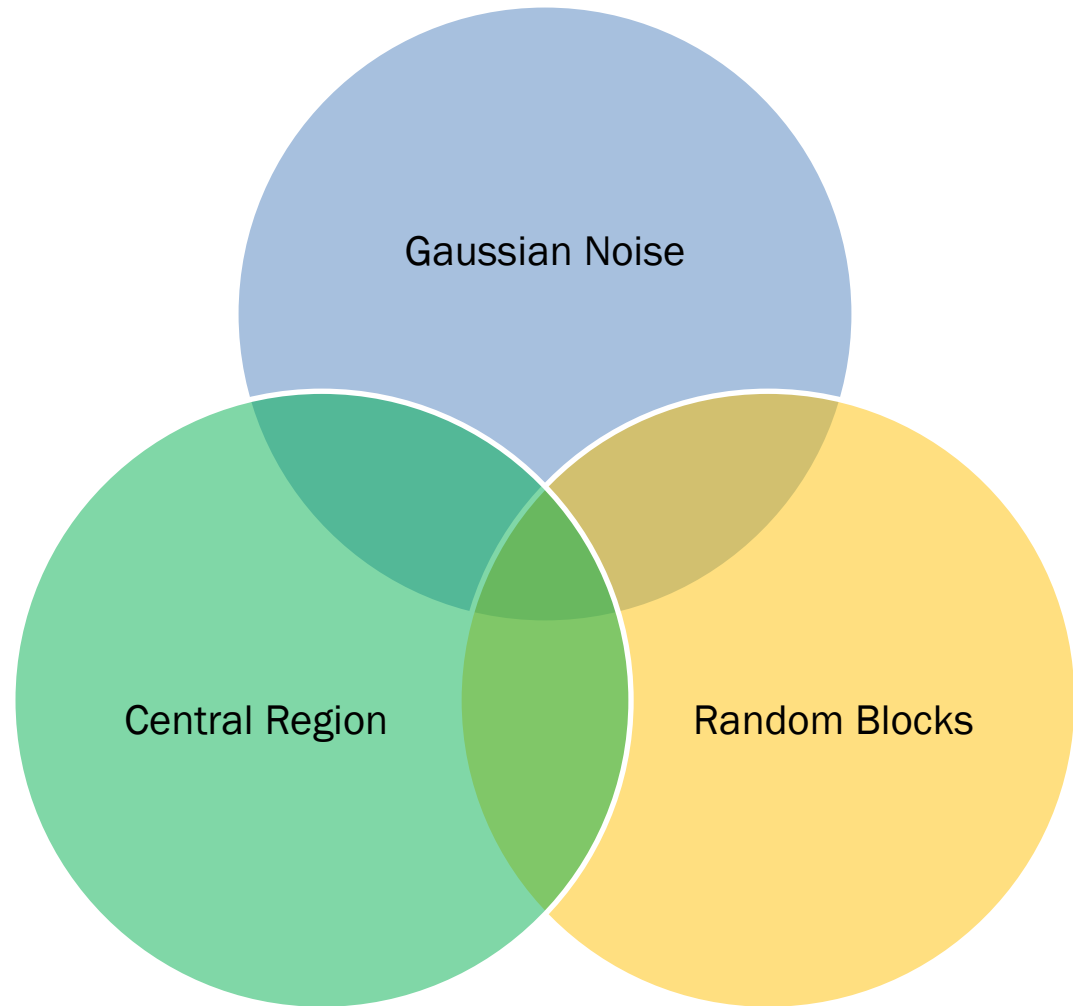
sigma=3



sigma=7



## 2. Advanced transformations



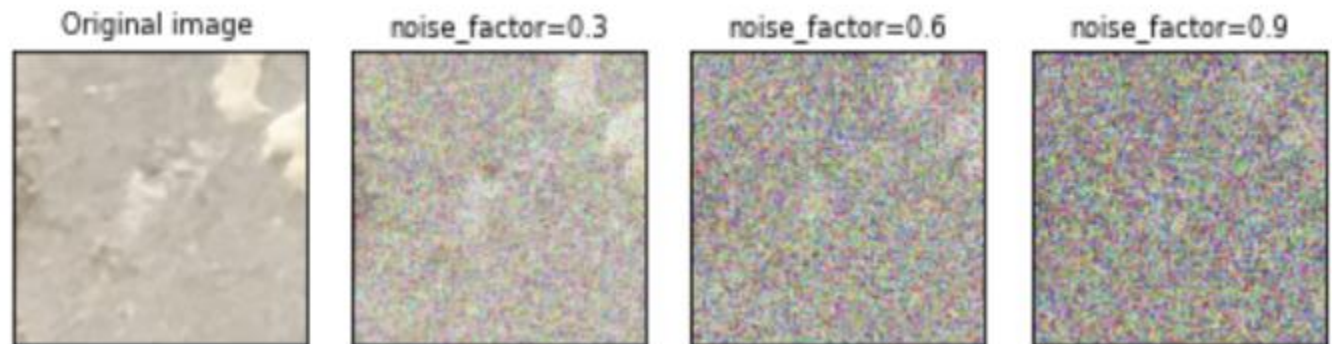


# Gaussian Noise

The Gaussian Noise is a popular way to add noise to the whole dataset, forcing the model to learn the most important information contained in the data.

```
def add_noise(inputs, noise_factor=0.3):  
    noisy = inputs + torch.randn_like(inputs) * noise_factor  
    noisy = torch.clip(noisy, 0., 1.)  
    return noisy
```

```
noise_imgs = [add_noise(T.ToTensor()(orig_img), noise_factor) for noise_factor in (0.3, 0.6, 0.9)]  
noise_imgs = [T.ToPILImage()(noise_img) for noise_img in noise_imgs]  
plot(noise_imgs, col_title=["noise_factor=0.3", "noise_factor=0.6", "noise_factor=0.9"])
```



# Random Blocks

Square patches are applied as masks in the image randomly. The higher the number of these patches, the more the neural network will find challenging the problem to solve.

```
def add_random_boxes(img, n_k, size=32):  
    h, w = size, size  
    img = np.asarray(img)  
    img_size = img.shape[1]  
    boxes = []  
    for k in range(n_k):  
        y, x = np.random.randint(0, img_size-w, (2,))  
        img[y:y+h, x:x+w] = 0  
        boxes.append((x, y, h, w))  
    img = Image.fromarray(img.astype('uint8'), 'RGB')  
    return img
```

```
blocks_imgs = [add_random_boxes(orig_img, n_k=i) for i in (10, 20)]  
plot(blocks_imgs, col_title=["10 black boxes", "20 black boxes"])
```

Original image



10 black boxes

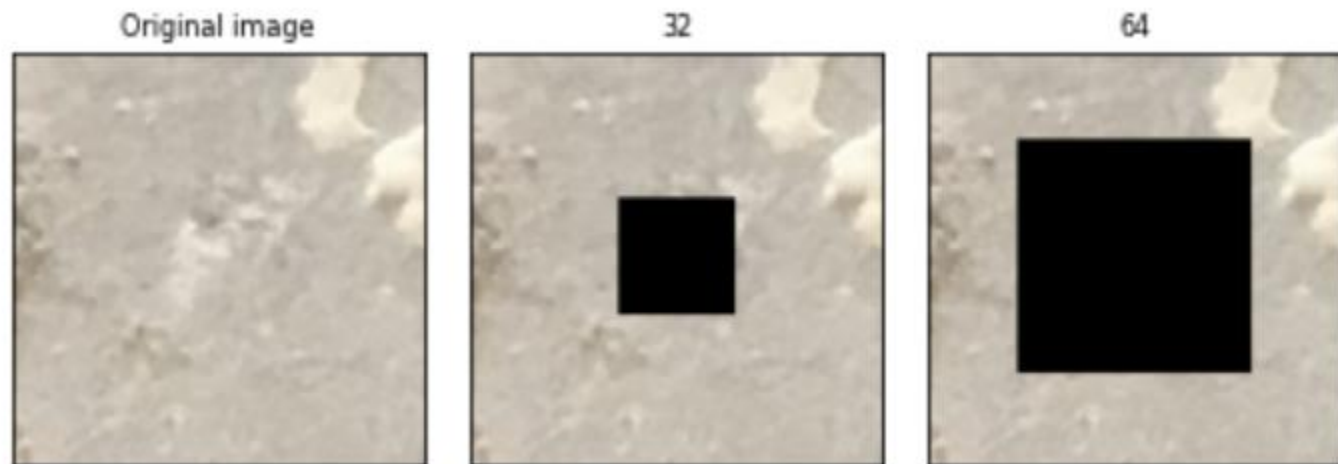


20 black boxes



# Central Region

```
def add_central_region(img, size=32):  
    h, w = size, size  
    img = np.asarray(img)  
    img_size = img.shape[1]  
    img[int(img_size/2-h):int(img_size/2+h), int(img_size/2-w):int(img_size/2+w)] = 0  
    img = Image.fromarray(img.astype('uint8'), 'RGB')  
    return img  
  
central_imgs = [add_central_region(orig_img, size=s) for s in (32, 64)]  
plot(central_imgs, col_title=["32", "64"])
```





# **ALBUMENTATIONS**



# Why Albumentations

- Albumentations **supports all common computer vision tasks** such as classification, semantic segmentation, instance segmentation, object detection, and pose estimation.
- The library provides **a simple unified API** to work with all data types: images (RGB-images, grayscale images, multispectral images), segmentation masks, bounding boxes, and key points.
- The library contains **more than 70 different augmentations** to generate new training samples from the existing data.
- Albumentations is **fast**. We benchmark each new release to ensure that augmentations provide maximum speed.
- It **works with popular deep learning frameworks** such as PyTorch and TensorFlow. By the way, Albumentations is a part of the PyTorch ecosystem.
- **Written by experts**. The authors have experience both working on production computer vision systems and participating in competitive machine learning. Many core team members are Kaggle Masters and Grandmasters.
- The library is **widely used** in industry, deep learning research, machine learning competitions, and open source projects.



# Example

Original image



RGBShift



HueSaturationValue



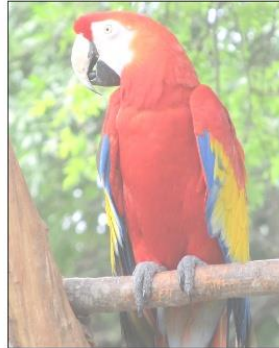
ChannelShuffle



CLAHE



RandomContrast



RandomGamma



RandomBrightness



Blur



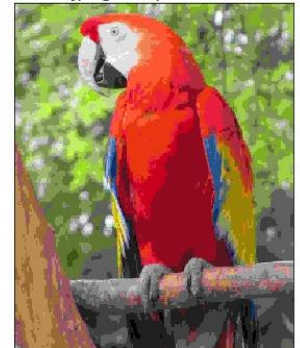
MedianBlur



ToGray



JpegCompression





# Installation

Install the latest stable version from PyPI

```
pip install -U albumentations
```

Install the latest version from the master branch on GitHub

```
pip install -U git+https://github.com/albumentations-team/albumentations
```

# Image Augmentation with Albumentations

```
from PIL import Image
import numpy as np
import albumentations as A
import matplotlib.pyplot as plt

orig_img = Image.open("img1.jpg")
image = np.array(orig_img )
plt.imshow(image)
plt.axis('off')
plt.show()
```



(256, 256, 4)



# Resize and Crop

```
transform_resize = A.Resize(width=64, height=64)
```

```
transform_cc = A.Compose([  
    A.Resize(width=128, height=128),  
    A.CenterCrop(width=32, height=32),  
])
```

```
transform_rc = A.Compose([  
    A.Resize(width=128, height=128),  
    A.RandomCrop(width=32, height=32),  
])
```

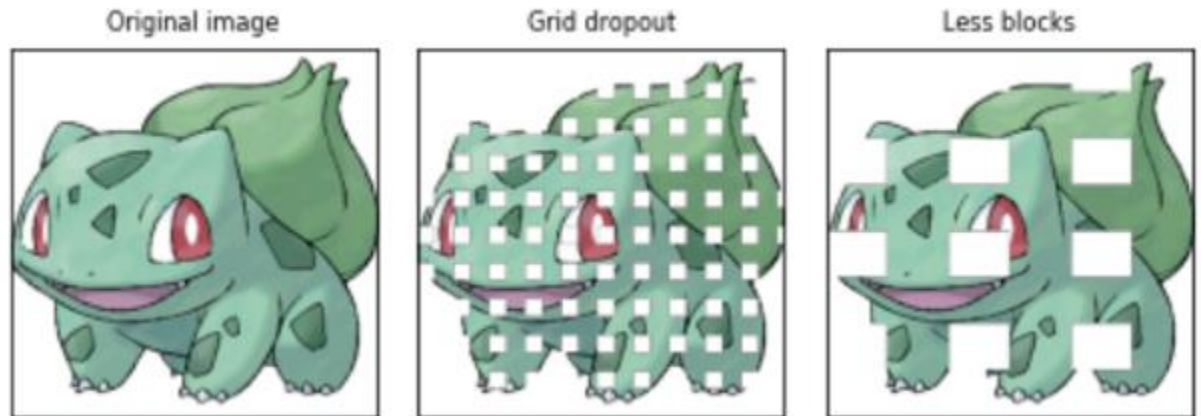
```
transformed_res = transform_resize(image=image)  
transformed_cc = transform_cc(image=image)  
transformed_rc = transform_rc(image=image)  
plot([transformed_res['image'],transformed_cc['image'],tr  
ansformed_rc['image']],col_title=["Resize 64x64","Resize  
& Center Crop","Resize & Random Crop"])
```



# Grid Dropout

```
transform_grid= A.GridDropout(p=1.0)
transformed_grid = transform_grid(image=image)
transform_grid2= A.GridDropout(p=1.0,holes_number_x=3,holes_number_y=4)
transformed_grid2 = transform_grid2(image=image)

plot([transformed_grid['image'],transformed_grid2['image']],col_title=["Grid dropout","Less blocks"])
```



# Different types of blur

```
transform_blur = A.Blur(p=1.0)
transform_mblur = A.MedianBlur(p=1.0)
transform_gblur = A.GaussianBlur(sigma_limit=9, p=1.0)

transformed_blur = transform_blur(image=image)
transformed_gblur = transform_gblur(image=image)
transformed_mblur = transform_mblur(image=image)
plot([transformed_blur['image'],transformed_gblur['image'],
transformed_mblur['image']],col_title=["blur","gaussian
blur","median blur"])
```



# Gaussian noise

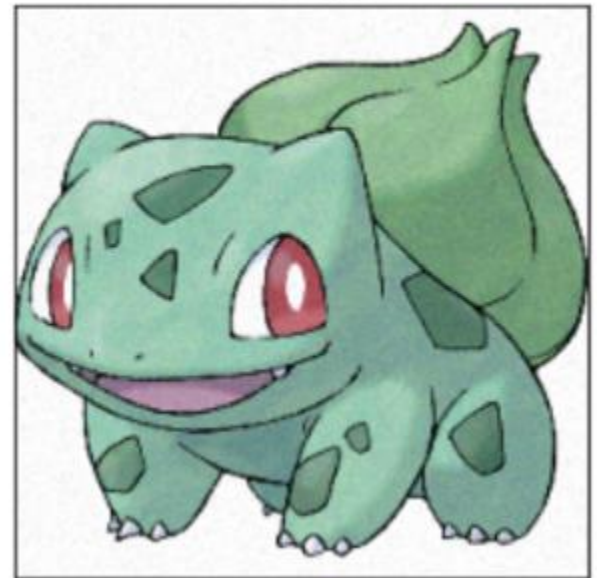
```
transform = A.GaussNoise(var_limit=350.0, p=1.0)
```

```
transformed_gnoise = transform(image=image)  
plot([transformed_gnoise['image']], col_title=["Gaussian N  
oise"])
```

Original image



Gaussian Noise





# Integration with Pytorch

```
from torchvision import transforms, datasets
from torch.utils.data import Dataset, DataLoader
from albumentations.pytorch import ToTensorV2
import os
import cv2
import random
import torchvision

dataset_directory = "../input/pokemon-images-  
dataset/pokemon/pokemon"
pokemon_filepaths = sorted([os.path.join(dataset_directory, f  
) for f in os.listdir(dataset_directory)])
correct_images_filepaths = [i for i in pokemon_filepaths if c  
v2.imread(i) is not None]


random.seed(42)
random.shuffle(correct_images_filepaths)
n = len(correct_images_filepaths)
n_train = int(n*0.8)
train_images_filepaths = correct_images_filepaths[:n_train]
test_images_filepaths = correct_images_filepaths[n_train:]
print(len(train_images_filepaths), len(test_images_filepaths)  
)
```



```
class PokemonDataset(Dataset):
    def __init__(self, images_filepaths, transform=None):
        self.images_filepaths = images_filepaths
        self.transform = transform

    def __len__(self):
        return len(self.images_filepaths)

    def __getitem__(self, idx):
        image_filepath = self.images_filepaths[idx]
        image = cv2.imread(image_filepath)
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
        if self.transform is not None:
            image = self.transform(image=image) ["image"]
        return image
```

A stack of five smooth, dark, rounded stones is arranged vertically on a reflective surface, likely water. The stones are dark in color, possibly black or dark grey, and their smooth texture is evident. The surface they rest on is calm, creating clear reflections of the stones. The background is a soft, out-of-focus light blue and white, suggesting a bright, open environment like a beach or a lake at dawn or dusk. The overall mood is serene and minimalist.

```
train_transform = A.Compose([
    A.Resize(height=128, width=128),
    A.Rotate(),
    A.GaussianBlur(sigma_limit=9, p=0.5),
    A.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5,
0.5)), ToTensorV2(), ])

test_transform = A.Compose([
    A.Resize(height=128, width=128),
    A.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5,
0.5)), ToTensorV2(), ])

train_dataset = PokemonDataset(images_filepaths=train_images_filepaths, transform=train_transform)
test_dataset = PokemonDataset(images_filepaths=test_images_filepaths, transform=test_transform)
train_loader = DataLoader(dataset=train_dataset, batch_size=16, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=16, shuffle=False)
```





```
def show_img(img):  
    plt.figure(figsize=(20,16))  
    img = img * 0.5 + 0.5  
    npimg = np.clip(img.numpy(), 0., 1.)  
    plt.imshow(np.transpose(npimg, (1, 2, 0)))  
    plt.show()  
  
data = iter(train_loader)  
images = data.next()  
show_img(torchvision.utils.make_grid(images))
```







Part A

# **STYLE TRANSFERRING**

# 1. What's Style Transferring

Style Transferring will use neural representation to separate, recombine content and style of arbitrary images, providing a neural algorithm for the creation of artistic images.

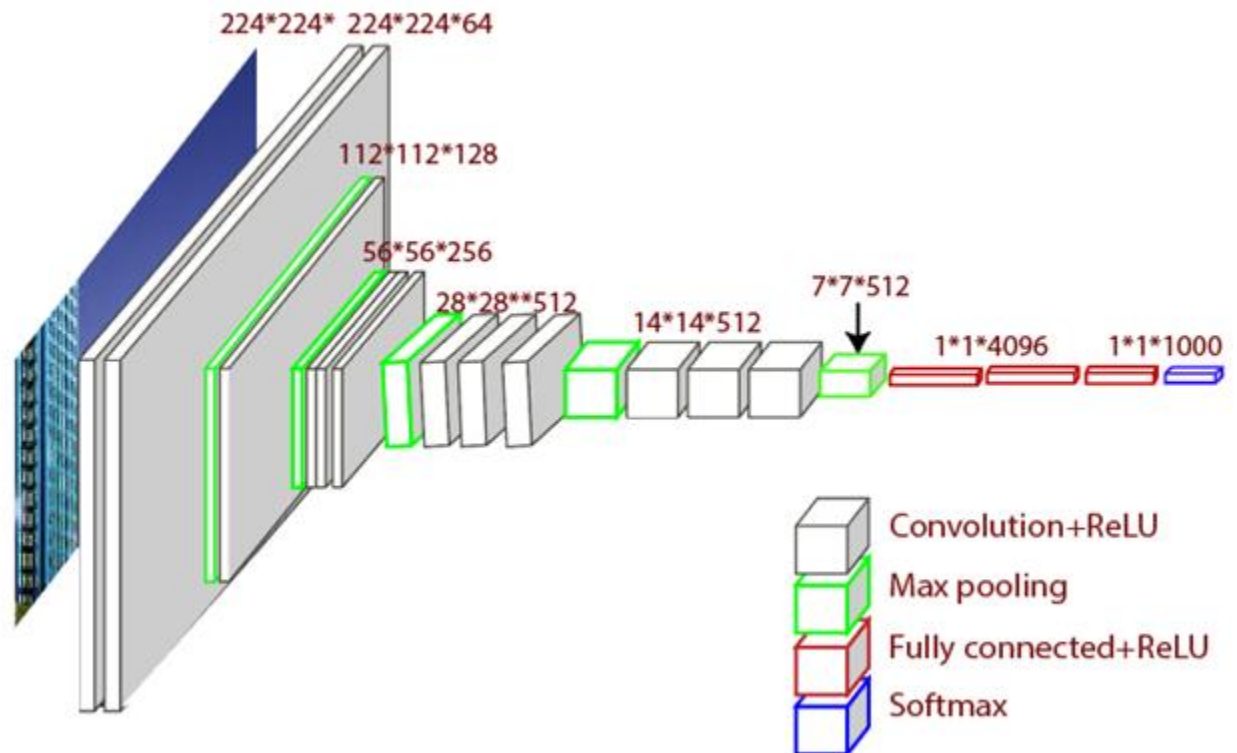
Neural style transfer is a way to generate images in the style of another image. The neural-style algorithm takes a content-image (a style image) as input and returns the content image as if it printed using the artistic style of the style image.

Style Transferring



# VGG-19 model

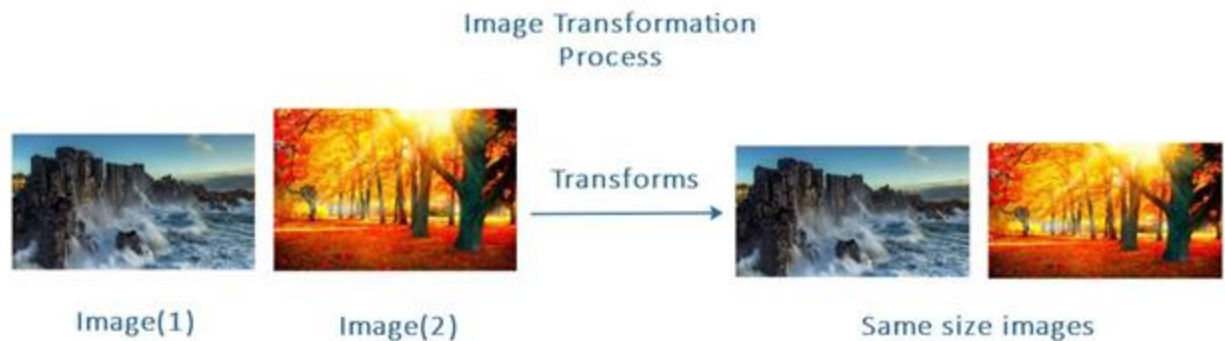
The VGG model was introduced by Simonyan and Zisserman. VGG-19 is trained on more than a million images from the ImageNet database. This model has 19 layers deep neural network, which can classify images into 1000 objects categories.





# Image loading and transformation

We have a content image, and style image and the target image will be the combination of both these images. Not every image needs to have the same size or pixel. To make the images equal, we will also apply the image transformation process.



# Feature Extraction



# Gram Matrix



# Optimization process







**Any Questions?**